Space Challenges Program 2026

Technical Design Document

Armonia

A Three-Axis Reaction-Wheel Balancing Cube

Project Team

Athanasia Nikolova, Andrea Liang, Deyan Nikolaev Vlaev

Niccolo Tonetto, Suvanna Viriyanti Wu, Pablo Urioste

Document version v1.0

Date of issue August 18, 2026

Contents

List of Figures

List of Tables

1 Introduction

1.1 Background & Problem Statement

The Cubli is a reaction-wheel inverted pendulum: a cube-shaped platform that balances on an edge or a corner using only internal actuation, with no external supports or contact forces beyond the single line or point of contact with the ground. The concept was introduced by the Institute for Dynamic Systems and Control (IDSC) at ETH Zurich [1, 2], whose original prototype (15 × 15 × 15 cm) mounts one brushless DC motor and flywheel on each of three orthogonal faces. Commanding a change in wheel speed produces, by conservation of angular momentum, an equal and opposite reaction torque on the cube body; braking a wheel abruptly transfers its stored angular momentum to the body in a near-impulsive event, which is what allows the cube to *jump up* from resting on a face into a balance on an edge or corner.

Dynamically, the system is an *underactuated*, nonlinear, unstable-equilibrium control problem: three actuators (the wheels) must stabilise a body with more degrees of freedom than actuated inputs, about equilibria (edge and corner balance) that are open-loop unstable and increasingly coupled across axes as the contact geometry shrinks from a line (edge) to a point (corner). This is structurally the same problem as the classical inverted pendulum, extended from a single rotational axis to a fully three-dimensional, momentum-exchange-actuated body — which is why the Cubli has become a standard benchmark platform in the controls community over the past decade.

Alongside the technical problem sits an educational one. Attitude determination and control is taught almost everywhere as mathematics on a whiteboard, because the hardware it describes is unreachable: flight units are expensive, single-shot and locked behind qualification procedures, so a student may complete an entire aerospace curriculum without ever commanding a reaction wheel and watching a body respond. What is missing is not theory but a system a student is allowed to build, mis-wire, crash and fix. The Cubli is exactly that system — it runs the same momentum-exchange physics as a real satellite, on a desk, at a cost a school can absorb.

The problem this project addresses is therefore threefold: (i) derive and validate a dynamic model and controller capable of stabilising this underactuated system about its unstable equilibria and driving it through commanded reorientations; (ii) realise this in hardware as a modular, largely parametric design, since final actuator and sensor specifications are not yet confirmed at the time of writing; and (iii) do both in a form that high-school and university students can reproduce, so that the mechanical build, the electronics, the firmware and the control law are documented as *reasoned decisions* a learner can follow and re-derive, not as finished numbers to be copied. The target capability set, in order of increasing difficulty, is edge balance, corner balance, controlled rotation, jump-up, and walking (Section 1.3).

No single subsystem is exotic, but each sits close to a physical bound, and the bounds are coupled: relieving one tightens another. Four constraints drive the design, and sizing is therefore treated as the first engineering activity rather than a detail settled once parts arrive.

- **The jump-up has no energy margin.** The required per-wheel inertia scales as $ML^{3/2}$ (Section 4.1.1), so mass added anywhere demands a larger wheel, and the wheel is itself part of M. Sizing is a fixed-point iteration, not a one-pass calculation.
- **Braking torque fights the structure.** Momentum accumulated over seconds of spin-up is shed in tens of milliseconds, through one shaft, hub and motor mount — the parts the mass budget wants thin. The structural load case is set by the manoeuvre that proves the machine works.
- **Sensing, not the control law, sets the balancing limit.** Gyroscope bias drifts, and the accelerometer — the only absolute tilt reference — is corrupted by vibration the controller

itself generates. Wheel balance is therefore an estimation requirement, not a mechanical nicety.

- **The plant is nonlinear and constrained.** Corner balancing does not decouple into three single-axis problems: the axes are coupled through the contact point and the inertia tensor. The actuators saturate in both torque and speed, so stored momentum must be actively managed. Linearising about upright is adequate for balancing and inadequate for anything more [3].

1.2 Purpose, Application & Mission Context

The purpose of this project is educational: to make space engineering something a student can hold, build and debug. The intended application of the hardware described in this document is as a teaching platform for high-school and university students, through which a complete satellite attitude control chain — structure, electronics, firmware, estimation and control — is learned by constructing it rather than by reading about it. Every other claim in this section is subordinate to that one: the spacecraft analogy matters because it is what makes the exercise *real*, and the engineering rigour of the sections that follow matters because a teaching platform that does not actually work teaches nothing.

The Cubli is well suited to this role because it is not a simulation of a satellite subsystem; it is one, operated under gravity. The cube is therefore not presented here as a balancing novelty but as a terrestrial, gravity-loaded demonstrator of **momentum-exchange attitude control** — the same physical mechanism that governs how real spacecraft point themselves. Nearly every three-axis-stabilised satellite holds and changes its attitude using reaction wheels: flywheels whose commanded speed change produces a reaction torque on the spacecraft body, exactly as in the Cubli. Spacecraft typically carry three or four such wheels (often in a pyramidal arrangement for redundancy) — directly analogous to the Cubli's three orthogonal wheels — and this reaction-wheel approach is the standard for CubeSats and small satellites in particular, making the Cubli representative of a real, currently-flown hardware class rather than an idealized toy problem.

Reaction wheels are the actuator choice made here, and the alternative is worth stating explicitly because it is the standard smallsat comparison. Since this is a proof-of-concept demonstration of the control loop, a system actuated with orthogonal reaction wheels is well-suited to represent principles applicable to smallsat ADCS: the momentum-exchange actuation principle and closed-loop control architecture are the same as those used on real spacecraft, even though the Cubli's dynamics (balancing against gravity) differ from a satellite's free-fall attitude control problem. Reaction wheels are also structurally simpler and require less specialised hardware than control moment gyros, which offer higher torque density but at the cost of gimbal mechanisms and more complex control software — complexity that is typically only justified on larger spacecraft. For a low-cost, resource-limited project, a reaction-wheel Cubli therefore demonstrates the core actuation and control principles of smallsat ADCS in a practical and relevant way.

Framed as an Attitude Determination and Control System (ADCS) problem, the Cubli's core behaviours map directly onto a spacecraft's day-to-day pointing task:

- **Edge / corner balance** → regulation to a commanded reference attitude and disturbance rejection — holding a pointing setpoint against perturbations.
- **Controlled rotation** → a slew manoeuvre: reference tracking to a new commanded attitude, exactly as a satellite re-aims a payload or antenna.
- **Jump / walk** → momentum-exchange hopping mobility, flight-proven on small bodies where wheeled traction is impossible (Figure 1). Walking, as a sequence of controlled

directional hops, extends single-hop capabilities and represents an active research frontier (Section 1.3).

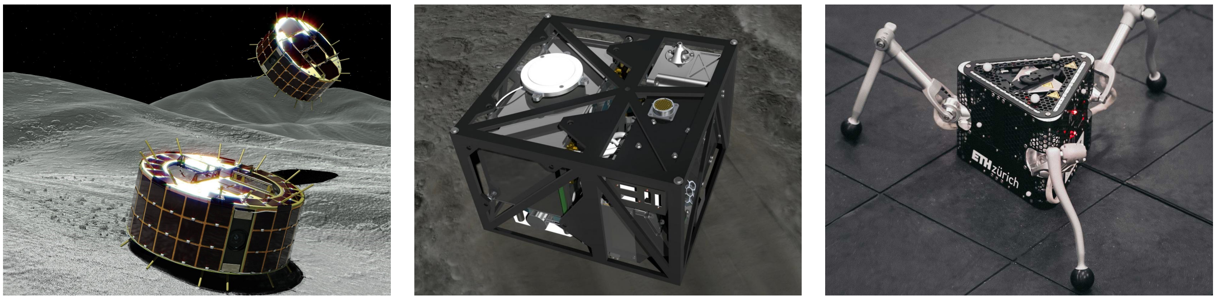

(a) MINERVA-II1 (b) MASCOT (c) SpaceHopper

Figure 1: Momentum-exchange actuation in flight and research hardware: (a) JAXA's MINERVA-II1 rovers [4] and (b) the MASCOT lander [5], deployed by Hayabusa2 onto asteroid Ryugu (2018), hop across the surface and self-right by spinning up and braking an internal flywheel or eccentric mass; (c) ETH Zurich's SpaceHopper [6] extends this principle to controlled, directional hopping in low gravity.

The subsystem exercised here is therefore the one a spacecraft uses for pointing: momentum-exchange attitude control under actuator saturation, with desaturation, on an underactuated body whose inertia is imperfectly known. Under gravity the destabilising torque never relents, so an inadequate controller falls over immediately and unambiguously — the demonstration is self-evaluating, which a free-fall attitude testbed is not. That property is what makes the platform teachable: the student needs no instructor and no test verdict to know whether the design is correct, because the cube either stands up or it does not.

Because the cube is a complete system rather than a single experiment, building it forces a student through the whole engineering chain, each discipline entered because the next one depends on it:

- **Design & CAD** → parametric modelling, tolerance and fit, centre-of-mass placement, and design for 3D printing (Section 5, Appendix B).
- **Physics & sizing** → conservation of angular momentum turned into numbers: wheel inertia, torque and speed budgets, and the fixed-point iteration between mass and actuator that makes engineering trade-offs concrete (Section 4).
- **Structures & building** → printing, assembly, fastening and balancing real parts, and discovering that the braking manoeuvre which proves the machine works is also the load case that breaks it.
- **Electronics & power** → motor drivers, battery and power budgeting, harnessing, grounding and the noise a switching motor injects into a sensor sitting centimetres away (Section 6, Appendix C).
- **Coding & embedded systems** → real-time firmware, fixed-rate control loops, sensor drivers, telemetry and debugging on hardware where a software bug becomes a falling cube.
- **ADCS, estimation & control** → state estimation and sensor fusion from a noisy IMU, feedback design, saturation and momentum management — the same problems, and largely the same equations, as a CubeSat's pointing loop (Section 7, Appendix D).
- **Systems engineering** → requirements traced to capabilities, budgets that must close, staged verification, and the experience of one subsystem's decision propagating into three others (Sections 2.2 and 3).

The nominal Cubli problem was solved a decade ago on precisely manufactured hardware [1, 2, 3]; reproducing it is an engineering exercise, and this document claims no more — but reproducibility is the point when the deliverable is a platform other students are meant to rebuild. A commodity build is adopted for four reasons. It can be crashed repeatedly, so the controller can be tested at its limits and a learner can fail without consequence; its loose tolerances, imperfectly balanced wheels and drifting IMU supply the parameter uncertainty that robust control methods are intended to tolerate; a fixed motor, print volume and budget constrain the sizing, since a marginal design cannot be rescued by specifying a larger component; and every part is a catalogue item or a printed file, so the cost of entry stays within reach of a school or student team rather than a laboratory. The third of these motivates the adjustable ballast scheme of Section 4.3.

1.3 Goal & Objectives

Goal. *To build a low-cost, reproducible reaction-wheel cube that teaches high-school and university students the engineering of a real satellite — design, structures, electronics, embedded coding and ADCS — by having them build, program and fly one, and in doing so make space engineering an accessible, hands-on discipline rather than an abstract one.*

The goal is deliberately stated as an outcome for the student rather than as a performance figure for the machine, because the machine is the teaching instrument and not the product. It decomposes into two sets of objectives that must both be met: **technical** objectives, which the hardware and controller must demonstrate, and **educational** objectives, which the platform and this document must deliver. The technical objectives are the demanding ones — a cube that cannot balance teaches nothing — and they are treated in the rest of this document with full engineering rigour, on the premise that the pedagogy is only as good as the hardware underneath it.

Technical objectives. Each imposes a required capability. The mapping below keeps every requirement traceable from motivation to capability, and is what the sizing and control work below must be shown to satisfy.

Table 1: Objective-to-requirement traceability.

Objective	Required capability
Edge balance	The system must actively maintain dynamic stability and regulate its posture about a single line-contact equilibrium position.
Corner balance	The system must maintain multi-axis closed-loop balance and attitude control around a highly unstable point-contact equilibrium.
Controlled rotation	The controller must accurately track dynamic reference trajectories to execute smooth slew maneuvers between distinct equilibrium states.
Jump-up (stretch)	The system must store sufficient wheel angular momentum and shed it near-impulsively through a braking event to tip itself from a face rest onto an edge or corner, and then capture the resulting attitude into balance.
Walking (stretch)	The system must execute a sequential series of dynamic jump-up maneuvers and rapidly re-stabilize itself upon each landing impact.

These map one-to-one onto the applications of Section 1.2: (1)–(2) demonstrate pointing regulation and disturbance rejection of increasing difficulty; (3) demonstrates slew/reference-tracking;

(4)–(5) demonstrate momentum-exchange mobility of the kind flown on Hayabusa2, first as a single hop and then as a sequence of them. Milestones (1)–(3) are treated as committed deliverables for this design; jump-up (4) and walking (5) are scoped as stretch objectives, both contingent on the jump-up feasibility check (angular momentum required to tip the cube vs. available motor/wheel torque) and on hardware confirmation, to be resolved once final components are known. Walking additionally depends on (4) being achieved, since each step is a jump-up followed by a re-capture.

The technical progression is also the teaching progression. Edge balance is reachable by a student who has understood a single-axis model and can tune one feedback law; corner balance demands the coupled three-axis treatment, a real estimator and honest inertia values; controlled rotation adds trajectory tracking; jump-up adds near-impulsive momentum transfer and the braking load case; walking adds repeated impact and contact. A learner therefore enters at the level they can handle and is pulled upward by the next milestone, using the same hardware throughout.

Educational objectives. These state what the student is expected to be able to do, and what the platform must provide for that to be possible.

Table 2: Educational objectives and what the platform must provide to deliver them.

Learning objective	Required provision
Build it	The mechanical design must be printable on hobby-grade equipment and assembled with catalogue fasteners and standard tools, so that construction is a student activity rather than a workshop service.
Wire it	The electrical architecture must be documented to schematic, pinout and harness level, with a bring-up procedure that isolates faults one subsystem at a time.
Program it	The firmware must expose the control loop, the estimator and the tunable gains as editable, readable code, with telemetry sufficient for a student to see why a run failed.
Understand it	Every sizing and control result must be derived from stated assumptions and reproducible from the equations given, so the numbers can be re-derived and challenged rather than trusted.
Break and fix it	The build must tolerate repeated falls and be repairable from spare printed parts and stocked components, so that experimentation carries no penalty beyond time.
Reproduce it	The design must be transferable: parametric CAD, a complete bill of materials and staged verification tests, so another school or team can rebuild the platform from this document.

The educational objectives are not a separate workstream layered on top of the engineering; they are constraints *on* it, and they have already changed design decisions. The commodity, crashable build of Section 1.2, the parametric CAD of Section 5, the adjustable ballast of Section 4.3 that keeps the cube tunable after assembly, the staged 1-DoF-before-3-DoF validation of Section 3 and the bill of materials of Appendix A all exist in the form they do because the platform has to be rebuildable and survivable in a student's hands, not only because they are sound engineering.

2 Project Organisation

2.1 Team Organisation

The project is developed by a multidisciplinary team, with responsibilities distributed across the main technical domains of the Cubli: mathematical modelling, control, embedded software, mechanical and three-dimensional design, electrical build, and system integration and validation. Table 3 states each member's role and the scope they are accountable for.

Table 3: Team members, contact details and primary roles.

Member	Contact	Role and accountable scope
Pablo Urioste Alarcón	pb.urioste4@gmail.com +34 654 491 174	**Project management and hardware-facing software.** Schedule, gate criteria and task allocation; motor-driver configuration, the flight-computer–driver link, the sensor drivers and their calibration, and the fixed-rate loop that carries them. Owns the hardware–software seam. Assembles and edits this document.
Deyan Nikolaev Vlaev	vlaev.dido@gmail.com +359 877 082 499	**Dynamics and technical documentation.** Equations of motion, linearisation and jump-up energy sizing; drafting and deepening of this document against the model.
Niccolo Tonetto	niccolo.tonetto@gmail.com +39 327 569 4498	**Control stack and build support.** Single owner of the control and estimation software — estimator, safety and mode logic, and the consolidated control release — written against the sensor interface Pablo maintains. Fabricates the distribution board and cube harness; leads the sacrificial brake and jump-up scope.
Andrea Liang	andrea.liang.2006@gmail.com +31 06 5796 2976	**Plant identification and control design.** Sizing calculations and the inertia and control budgets; sole owner of rig parameter measurement, from the instrumented measurement through to the validated plant model the balancing controllers are designed on.
Suvanna Viriyanti Wu	Viriyanti.wu@gmail.com +39 329 074 7550	**1-DoF model and mechanical build.** Builds and iterates the single-axis prototype and its testbench, fabricates the reaction wheels, and supports the 3D design work.
Athanasia Nikolova	nasia.nikolova@mtel.gr +30 697 232 4491	**3D design and commercialisation.** 3D CAD, frame design and mechanical enclosures; monitors the design against the frozen mass and interface constraints; prototype assembly and the commercial pitch.

2.2 Work Breakdown Structure

This plan runs from project start to the final presentation on 20 August 2026. The work is grouped into five streams — investigation, requirements and documentation; control and software; hardware design and build; three-dimensional design; and electrical and electronics — led by the owners assigned in Table 3. The investigation is kept to the essentials needed to substantiate the design: a focused literature review, the dynamics derivation, and the wheel sizing that follows from it. Table 4 breaks the streams into individual tasks, each with its scope, single accountable lead and dates.

Table 4: Work breakdown structure, broken into individual tasks with scope, lead and dates, grouped by workstream.

ID	Task	Scope	Lead	Dates
Investigation, Requirements & Documentation				
A1	Literature & state of the art	Review the ETH Cubli, the one-wheel Cubli, CubeSat ADCS practice and the hopping-mission precedents	Suvanna	23–25 Jul
A2	Problem statement & framing	Lock the ADCS application narrative that the rest of the document refers back to	Pablo	23–24 Jul
A3	Scope & requirements	Define scope and build the requirements table (functional, performance, design constraint) with verification methods	Pablo	24–25 Jul
A4	Milestones & roles	Prototype staging, binary gate criteria, and a one-page RACI of owners	Pablo	23–25 Jul
A5	Risk register	Trigger, impact, owner and mitigation for the leading risks	Pablo	28–29 Jul
A6	1-DoF equations of motion	Lagrangian derivation of the reaction-wheel pendulum with friction terms	Deyan	24–25 Jul
A7	3D equations of motion	Coupled three-wheel dynamics for the edge and corner pivots	Deyan	27–28 Jul
A8	Linearisation & controllability	Equilibria, linear models, controllability/observability checks	Deyan	29–31 Jul
A9	Jump-up sizing & budget	Energy-method wheel sizing and the parametric mass/inertia/CoM budget	Deyan	30 Jul–2 Aug
A10	TDD draft (all sections)	Every member drafts their own section; placeholder figures acceptable	Deyan	25–28 Jul
A11	TDD deepen + rig evidence	Modelling and control sections in full, with the rig evidence and measured parameters available at the time of writing	Deyan	29 Jul–2 Aug
A12	TDD v1 assemble & submit	Consistency pass, figures, references, submission of the first assessed deliverable	Pablo	3 Aug
A13	Assessor feedback triage	Triage the v1 assessment, assign each comment to an owner, and re-baseline roles and task allocation for v2	Pablo	4–7 Aug
A14	TDD v2 corrections & second draft	Address the assessor comments and fold in the build, bring-up and test results since v1. Ends 12 Aug: no work 13 Aug	Pablo	8–12 Aug
A15	TDD v2 final submission	Final consistency pass and submission of the corrected document	Pablo	14 Aug
Control & Software				
B1	Sim environment + plant	Simulation setup and the 1-DoF plant, validated by energy conservation	Andrea	23 Jul–3 Aug
B2	LQR design (1-DoF, sim)	Tune the regulator and sweep robustness to inertia error, delay and saturation	Andrea	26 Jul–6 Aug
B3	State estimator	Complementary tilt filter, tested against injected noise and bias	Niccolo	27 Jul–7 Aug
B3b	Software interface contract	Fix the sensor and command structs passed between the sensor loop and the controller, with the units and sign convention of every field	Pablo	3–4 Aug

ID	Task	Scope	Lead	Dates
B4	Flight computer & comms	Teensy 4.1 architecture, loop rates, driver configuration and the CAN-FD command/reply protocol	Pablo	24 Jul–7 Aug
B5	Sensor drivers + cal	IMU and encoder drivers, gyro/accel calibration and the lever-arm correction	Pablo	28 Jul–7 Aug
B6	Safety & mode logic	Watchdog, tilt-limit disarm, wheel-speed cap and the mode state machine	Niccolo	30 Jul–7 Aug
B7	1-DoF control design complete	Consolidated 1-DoF control stack — plant, regulator, estimator and safety logic — released as the rig baseline	Niccolo	8–9 Aug
B8	System identification	Measure K_t, wheel inertia, friction and loop latency on the rig. Single owner with E6	Andrea	7–9 Aug
B9	3D model + LQR (sim)	3D plant and the edge/corner regulators with the CAD inertia tensor	Andrea	10–14 Aug
B10	EKF / MEKF estimator	3D estimator with gyro-bias states; characterise the yaw drift	Niccolo	12–15 Aug
B11	Edge balance (S2a)	Retune the edge controller on the assembled cube; log disturbance recovery	Andrea	15–17 Aug
B12	Corner balance + slew (S2b/c)	Coupled three-axis corner balance and a commanded slew; report ADCS metrics	Andrea	17–19 Aug
B13	Jump-up (S3)	Spin-up profile and brake release for a face-to-edge jump. Sacrificial scope: yields to G6 and G7 if the schedule is contested	Niccolo	19–20 Aug
Hardware Design & Build				
C1	Envelope & layout study	Packaging of wheels, drivers, battery and electronics with a central CoM	Nasia	23–24 Jul
C2	Reaction-wheel design	Wheel to the inertia target: radius, rim, hub bolting to the motor bell	Suvanna	24–26 Jul
C3	Mounts & brackets	Motor mount, ring-magnet carrier and adjustable encoder bracket	Nasia	27–29 Jul
C4	1-DoF rig CAD + jigs	Single-axis rig and the balancing/assembly tooling	Suvanna	25–28 Jul
C6	Procurement order (round 1)	Order the long-lead items (rims, frame stock, connectors, spares)	Pablo	23–24 Jul
C7	Bench PSU & bring-up	Current-limited supply, safety kit, and single-axis electrical bring-up	Pablo	25–27 Jul
C8	Fabricate wheel rings	Print the PET-CF rings and hubs to the frozen wheel geometry	Suvanna	27 Jul–2 Aug
C8b	Wheel ballast & completion	Install the M6 bolt/nut ballast stations to the sizing specification and complete each wheel	Suvanna	4–6 Aug
C9	Balance & spin-test	Static-balance and containment spin-test every wheel, including a per-wheel ballast-count verification	Suvanna	5–6 Aug
C10	Assemble 1-DoF rig	Motor, balanced wheel, encoder gap, IMU and hard stops, on reinforced rig supports	Pablo	31 Jul–2 Aug
C11	Print structure + frame	Structural print set and the fabricated, orthogonal frame	Suvanna	8–12 Aug

ID	Task	Scope	Lead	Dates
C12	Perfboard + harness	Power/signal board and the full power and CAN harness	Pablo	8–12 Aug
C13	Cube integration + power-on	Assemble the three axes, then bench and battery smoke test	Pablo	13–15 Aug
C14	Brake mechanism build	Fabricate and fit the three servo brakes; bench-test the deceleration	Suvanna	15–18 Aug
3D Design				
D1	Concept sizing & fit study	Confirm that every component fits the 150 mm envelope and lay out the internal structure	Nasia	27 Jul–2 Aug
D2	Frame design freeze	Freeze the inertia-bearing frame geometry. Later features that do not affect the mass properties may still be added	Nasia	2 Aug
D3	Mass properties & budget closure	Export the inertia tensor and CoM and close the mass budget against the sizing analysis	Nasia	3–4 Aug
D4	Detail CAD	Brackets, electronics bay, harness routing and the electrical keep-outs to the frozen board outline	Suvanna	3–5 Aug
D5	Procurement order (round 2)	Consolidate the bill of materials arising from the 3D design and place the second order	Pablo	5 Aug
D6	Tolerance & clearance analysis	Fit and tolerance study, including the wheel sweep envelope and fastener access	Nasia	5–6 Aug
D7	Manufacturing package release	Issue the drawings, print set and assembly instructions for fabrication	Nasia	7 Aug
Electrical & Electronics				
E1	Wiring & placement design	Electrical architecture, wiring scheme and component placement for the single-axis rig, issued as the build reference	Pablo	28–31 Jul
E2	Rail, driver & encoder bring-up	5 V rail soldered and trimmed to 5.00 V under load; driver phase-wired and FOC-calibrated; MA600 encoder mounted at ≤ 0.5 mm air gap with verified read-back. Open-loop operation on the bench supply	Pablo	1–3 Aug
E3	Rig bus & sensor integration	Two-node CAN-FD bus with split termination at both physical ends, Teensy and IMU integrated on the rig harness	Pablo	4–5 Aug
E4	Board outline & floorplan	Driver placement from the encoder cable limit, CAN routing, and the perfboard floorplan issued as a frozen outline to CAD	Pablo	4–5 Aug
E5	1-DoF closed-loop bring-up	Closed loop on the rig with the disarm path and watchdog verified before any balance attempt	Pablo	6–7 Aug
E6	Rig parameter measurement	Instrumented measurement of K_t, wheel inertia, friction and loop latency. Merged with B8: one owner measures and models	Andrea	7–9 Aug

ID	Task	Scope	Lead	Dates
E6b	Mechanical & performance validation (1-DoF)	Reconcile the measured torque, inertia, friction and achievable momentum against the jump-up sizing (A9). Deviation beyond tolerance forces a re-sizing decision that day, while the wheel geometry can still change	Andrea	9 Aug
E7	Perfboard + battery power chain	Populate the distribution perfboard (two ground domains, one star point, split 5 V rail); XT60 chain with anti-spark, fuse and ≥ 50 V bulk capacitance; replicate the remaining actuation sets	Niccolo	8–11 Aug
E8	4-node bus + cube harness	Bus scaled to four nodes with stubs < 10 cm and an error soak at 5 Mbit/s; full power and signal harness fabricated and routed; Wi-Fi telemetry for untethered operation	Niccolo	11–12 Aug
E9	Cube power-on, 48 A + thermal	First battery power-on and the three-axis 48 A transient; driver thermals in the enclosed volume; all axes addressable	Pablo	14–15 Aug
E10	Brake servo rail + measurement	Servo on an isolated 5 V rail with PWM drive and flyback; rail stability under stall, brake closure time and spin-down brake torque measured. Sacrificial with C14/B13	Niccolo	16–19 Aug
Deliverables				
F1	Results, demo & final presentation	Slow-motion capture, results plots, deck, and the final presentation	Andrea	18–20 Aug

2.3 Schedule and Contingency Plan

The full-page schedule in Figure 2 places every task of Section 2.2 on the 23 July–20 August 2026 calendar, and Table 5 states the gate that closes each phase.

Structure of the plan. The first two weeks establish the analytical and design foundation: the dynamics derivation and wheel sizing on the documentation side, the single-axis rig and its wiring design on the hardware and electrical sides, and the simulation environment on the control side. These converge on 2 August, which carries two objectives: the first powered bring-up of the single-axis rig, taking the 5 V rail, the driver and the encoder to verified open-loop operation (M1); and the freeze of the three-dimensional frame geometry, which closes the mass and inertia budget against the sizing analysis (M2). The frame freeze releases the second procurement round on 5 August (M3), which is the last point at which parts can be ordered with useful delivery margin before cube assembly begins on 13 August.

The documentation stream carries two assessed submissions. The first is delivered on 3 August (G4) with the evidence available at that date; the assessor's comments are triaged during the following week and the corrected second draft is submitted on 14 August (M8), by which point the rig results and the integrated cube can both be reported.

Critical path and contingency. The critical path runs from the 2 August bring-up through the closed-loop rig (M4), the measured plant parameters (M6) and the single-axis control design (M7), into the three-dimensional controller and the balancing demonstrations. The single-axis rig is the protected item, since its balance result is the experimental evidence in the document. The three-dimensional design proceeds in parallel on the mechanical side, so that fabrication is not held by the control work.

The plan carries the full scope through to the jump-up demonstration. The final week is nonetheless tightly constrained: edge balance begins two days after cube assembly completes, and the jump-up gate falls on the day of the final presentation. The jump-up (S3) is therefore designated sacrificial scope by standing decision. If the balancing demonstrations require the time, the jump-up yields first, and the demonstration presented on 20 August is the corner balance and commanded slew. The brake mechanism build and its associated electrical work are sequenced after the edge-balance gate for the same reason.

Table 5: Milestone gates, from kick-off to the final presentation, with the criterion that closes each one.

Gate	Pass criterion	Date
G1	Long-lead order placed; requirements and plan locked	24 Jul
G2	Dynamics, linearisation and wheel sizing complete	31 Jul
M1	1-DoF rig bring-up: 5 V rail trimmed under load, driver commutating, encoder verified at ≤ 0.5 mm gap	2 Aug
M2	3D frame geometry frozen at 150 mm; mass, inertia and CoM budget closed against the sizing analysis	2 Aug
G4	TDD v1 delivered for assessment	3 Aug
M3	Procurement round 2 consolidated and ordered	5 Aug
M4	1-DoF closed loop live; disarm path and watchdog verified	7 Aug
M5	Manufacturing package released for fabrication	7 Aug
M6	Plant parameters measured on the rig and validated against simulation; measured torque, inertia and friction reconciled with the sizing analysis, or a re-sizing decision recorded	9 Aug
M7	1-DoF control design complete: plant, regulator, estimator and safety logic	9 Aug
G3	1-DoF rig balances $\geq$60 s unassisted	12 Aug
M8	TDD v2 corrected and finally submitted	14 Aug
G5	Cube integrated and powered; all axes addressable	15 Aug
G6	Edge balance (S2a) for $\geq$60 s with disturbance recovery	17 Aug
G7	Corner balance and commanded slew (S2b/c)	19 Aug
G8	Face-to-edge jump-up (S3) captured on video (sacrificial scope)	20 Aug
G9	Final presentation and demonstration	20 Aug

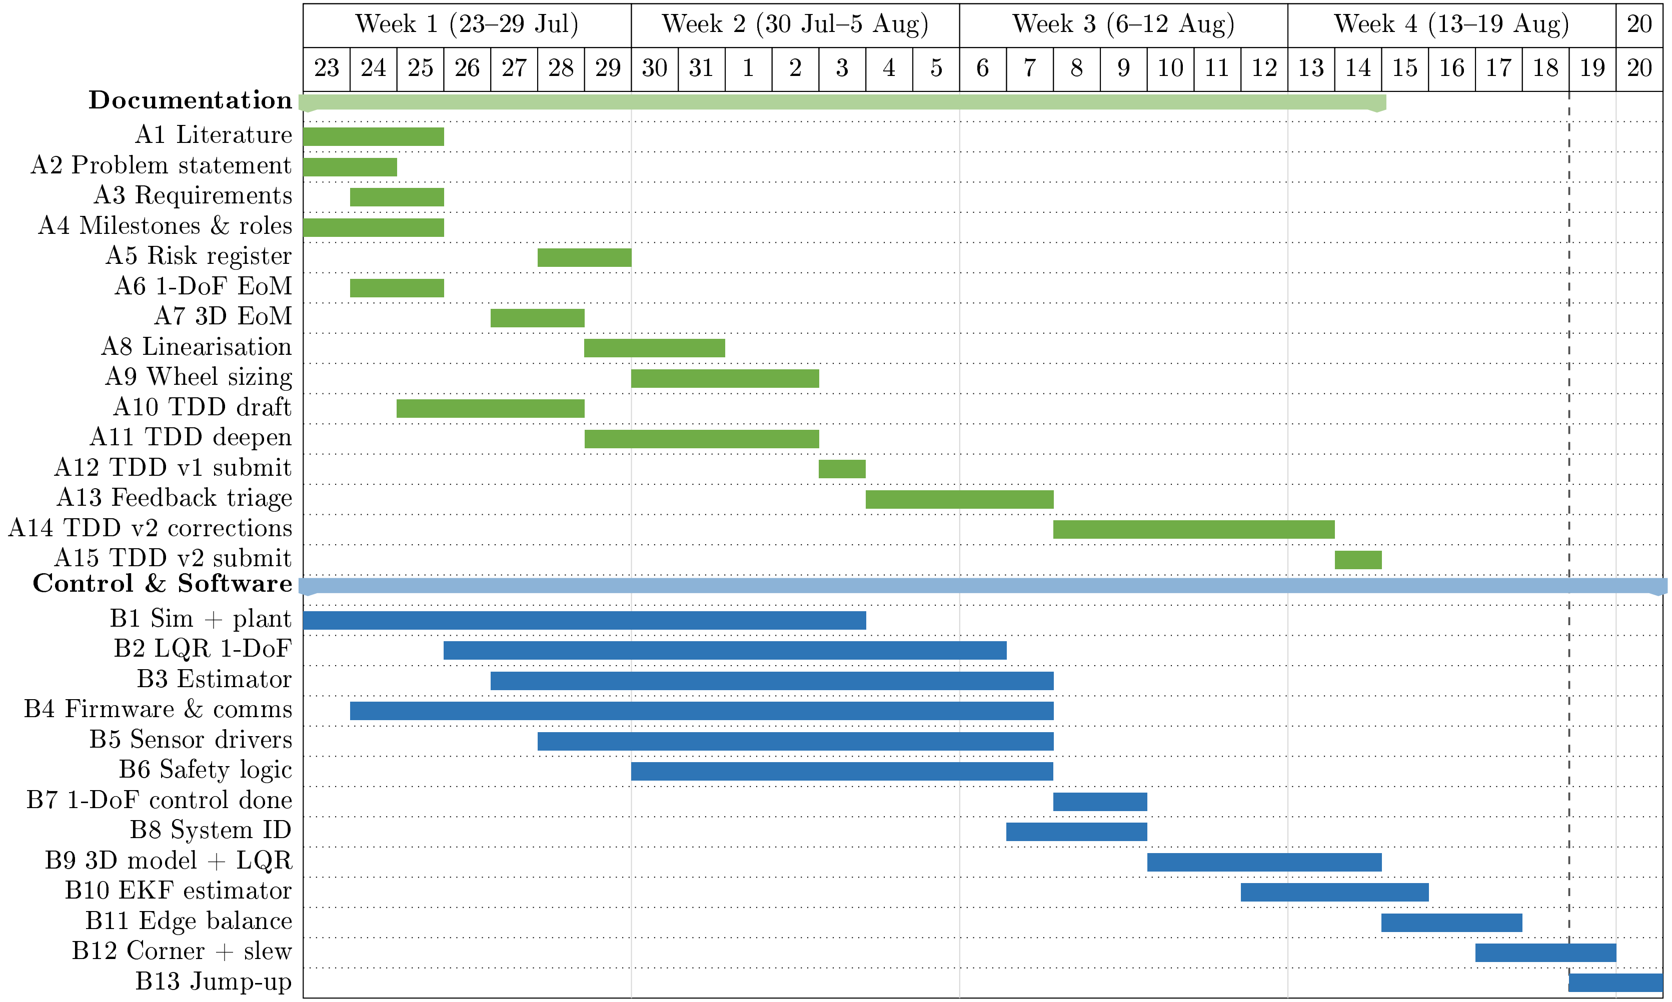

Figure 2: Project schedule, 23 July–20 August 2026 (sheet 1 of 3): documentation and control/software streams. The documentation stream (green) carries two assessed submissions, on 3 and 14 August. Control and software (blue) runs from the simulation environment through to the balancing demonstrations, with the three-dimensional extension concentrated in the final two weeks. The build streams continue in Figure 3 and the milestone gates in Figure 4.

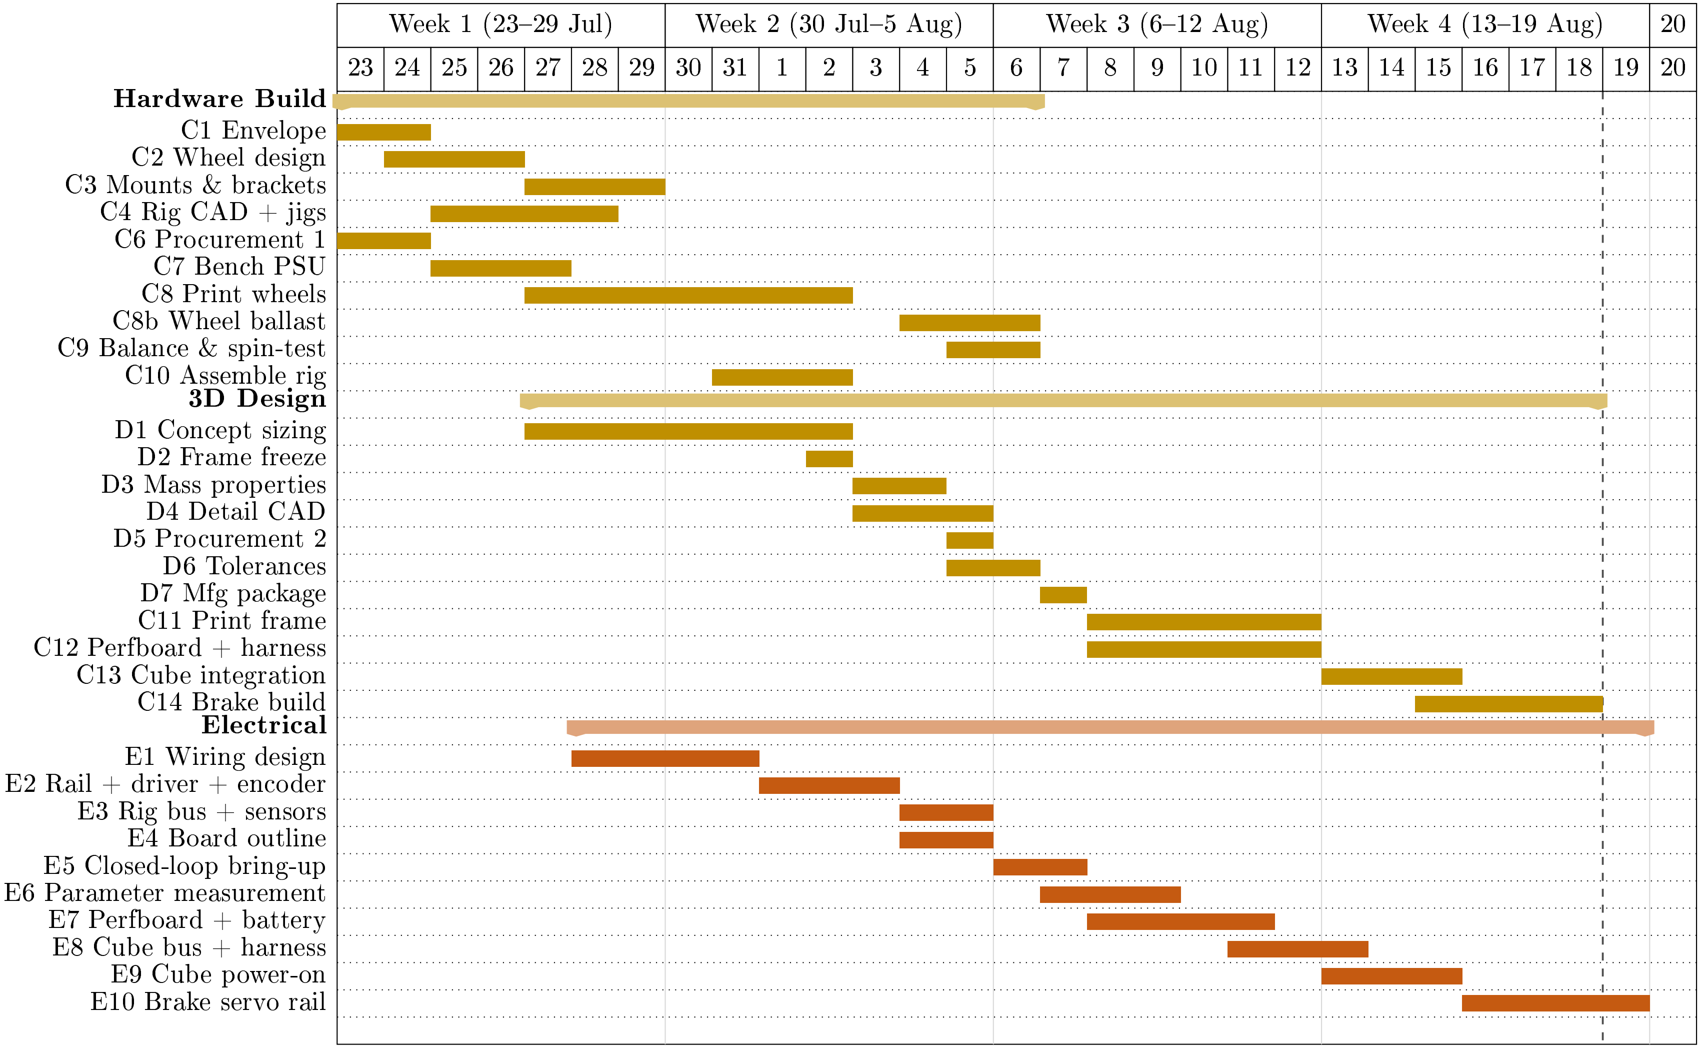

Figure 3: Project schedule, 23 July–20 August 2026 (sheet 2 of 3): hardware, three-dimensional design, electrical and deliverable streams, continued from Figure 2. Hardware (gold) separates the single-axis build from the three-dimensional design, and the electrical stream (orange) proceeds from the single-axis bring-up to the integrated cube. The milestone gates that close each phase are shown in Figure 4.

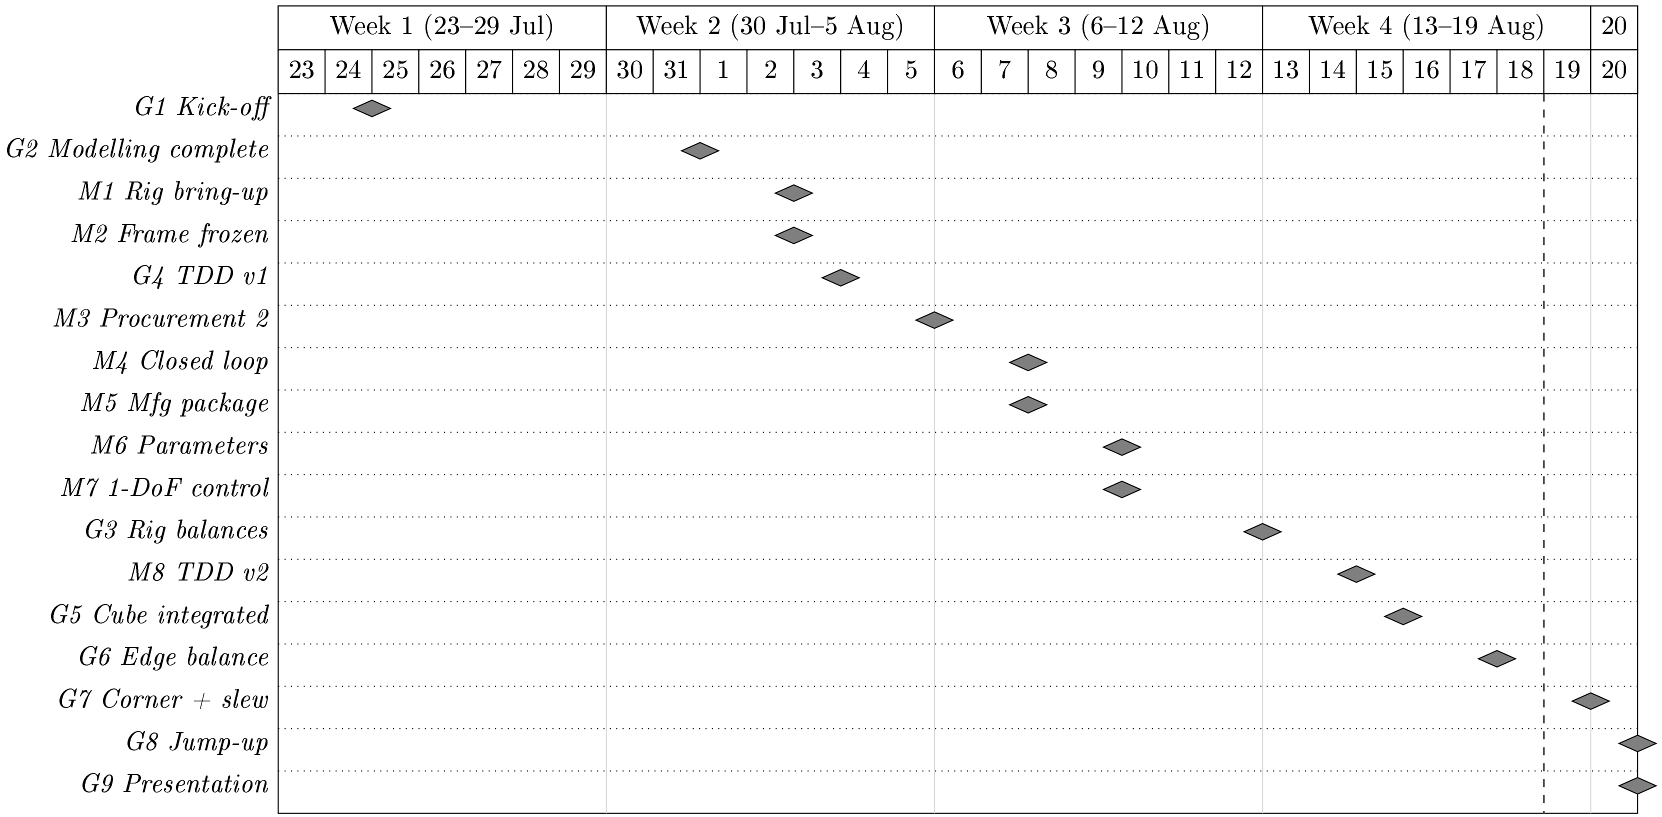

Figure 4: Project schedule, 23 July–20 August 2026 (sheet 3 of 3): the milestone gates that close each phase, continued from Figures 2 and 3. Grey diamonds mark the gates listed in Table 5, placed on the day each gate falls due.

3 Methodology

Section 2.2 assigns every task an owner and a date. This section states the reasoning that assignment is built on: the procedure the team follows to take the Cubli from a sizing target to a validated prototype, why its six stages run in this order and not another, and why the result is an iteration rather than a single pass through them. The Cubli is an interdisciplinary system whose subsystems cannot be sized, modelled or validated in isolation, which is why the methodology below is organised as a sequence of stages and validation gates rather than a single design pass.

3.1 Development and Iteration Workflow

Figure 5 is that procedure: six stages, each committing an output the next stage cannot start without, and three feedback paths that reopen an earlier stage when a downstream finding invalidates an upstream assumption.

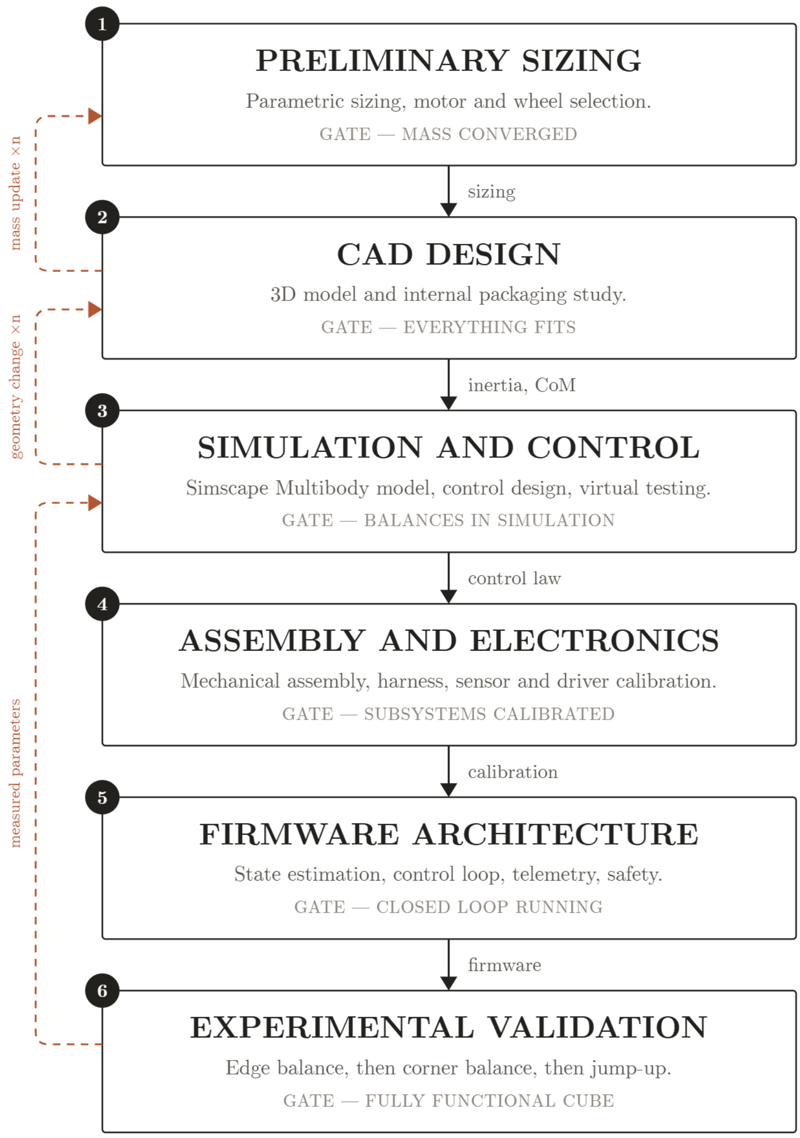

Figure 5: Development and iteration workflow, with the three feedback loops that keep sizing, CAD and control consistent with each other and with measured hardware.

Workflow steps Each stage sits where it does because the stage after it consumes an output only it produces:

1. **Preliminary Sizing** is the first stage of the project: it derives the physics behind the Cubli and justifies the design parameters and performance targets the rest of the design must meet.
2. **CAD Design** turns those targets into a 3D model and an internal packaging study, returning a measured inertia tensor and centre of mass where Stage 1 could only estimate them (Section 5).
3. **Simulation and Control** tests control feasibility against the CAD-derived inertia, mass and geometry rather than the Stage 1 estimate (Sections 3.2–3.4, 7).
4. **Assembly and Electronics** begins once the design is closed: the cube is assembled, the electronics are placed and wired, and every axis is tested before it is commanded (Section 6).
5. **Firmware Architecture** is developed in parallel with the previous stage: it integrates the electronics around the MCU and implements the control-dynamics loop that converts IMU orientation data into wheel-torque commands. Alongside the flight code, sub-programs are developed to test and calibrate each electronic module individually (Section 7.6.2).
6. **Experimental Validation** is the last stages: once integration concludes, results are tested and validated. The workflow accommodates iteration at every stage rather than assuming a single pass succeeds.

Iterative loop A single pass through that order is not sufficient, because three of its outputs correct an assumption made earlier in the same pass rather than only feeding forward:

- *Mass update.* CAD's as-modelled mass and CoM close the loop into Stage 1's assumed M (Section 4.2): the sizing reported there is the fixed point of this loop, not its first iteration.
- *Geometry change.* A control or simulation finding — inertia short of target, an infeasible slew, an interference the 2D model could not see — returns to CAD rather than being absorbed by re-tuning gains, since no gain fixes a wheel that does not fit or a body the controller cannot actuate against.
- *Measured parameters.* Once hardware exists, the torque constant, friction and latency measured on it (Section 3.3) replace the CAD/datasheet estimates Stage 3 was designed on, and the control gains are re-derived against the measured plant rather than flown on assumed values.

This is also why the schedule of Section 2.3 gates CAD on the sizing numbers rather than running the two in parallel: a frame frozen ahead of the inertia tensor is a frame that may have to be reprinted, which costs more than the days spent waiting on the number.

The full Cubli is underactuated and simultaneously unstable about all three axes when balanced on a corner (Section 1.3). Stage 3 of the procedure above is therefore itself incremental: a first, single-face 1-DoF prototype is built and modelled in two dimensions, so that the dynamics, the parameter identification procedure and the estimator can be validated on a plant simple enough to reason about directly; the same modelling, identification and verification workflow is then generalised to the second and final prototype, the fully three-dimensional cube, in Section 3.4. Each subsection below states the equations that define the problem at that stage and the method used to determine every parameter that appears in them; the control law itself — how those equations are turned into a feedback command — is derived separately in Section 7, since sizing, dynamics and control are three passes over the same plant rather than three independent problems.

3.2 2D Reaction-Wheel Pendulum

The 1-DoF prototype is a 3D-printed plate sized to one cube face, with a motor-driven momentum wheel at its centre and a braking mechanism at one corner. The plate pivots on a bearing at that corner, constraining motion to a single rotational degree of freedom in a vertical plane.

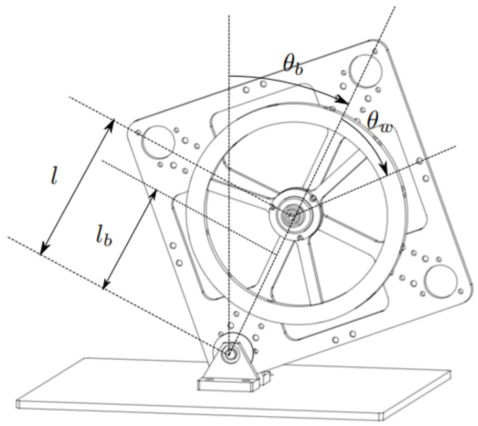

Figure 6: Frames and angles of the 2D reaction-wheel pendulum. θ_b is the plate's tilt from vertical about the pivot; θ_w is the wheel's rotation relative to the plate; l and l_b are, respectively, the pivot-to-motor-axis and pivot-to-plate-CoM distances along the plate's diagonal. Symbols per Table 6.

3.2.1 Equations of Motion

The nonlinear equations of motion are

$$\ddot{\theta}_b = \frac{(m_b l_b + m_w l)\, g \sin\theta_b - T_m - C_b\dot{\theta}_b + C_w\dot{\theta}_w}{I_b + m_w l^2}, \tag{1}$$

$$\ddot{\theta}_w = \frac{(I_b + I_w + m_w l^2)(T_m - C_w\dot{\theta}_w)}{I_w(I_b + m_w l^2)} - \frac{(m_b l_b + m_w l)\, g \sin\theta_b - C_b\dot{\theta}_b}{I_b + m_w l^2}, \tag{2}$$

with motor current-torque relation $T_m = K_m u$. Symbols are defined in Table 6 and used throughout this section.

Table 6: 2D pendulum model variables.

Symbol	Description	Unit
θ_b	Body tilt angle	rad
θ_w	Wheel rotation relative to body	rad
m_b, m_w	Body, wheel mass	kg
I_b, I_w	Body (about pivot), wheel (about motor axis) inertia	$\mathrm{kg\,m^2}$
l	Motor axis to pivot distance	m
l_b	Body CoM to pivot distance	m
g	Gravitational acceleration	$\mathrm{m/s^2}$
T_m	Motor torque	N m
C_b, C_w	Body, wheel damping coefficient	$\mathrm{kg\,m^2/s}$
K_m	Motor torque constant	N m/A
u	Motor current input	A

3.3 Parameter Identification & the 2D Testbed

CAD mass-properties simulation, accounting for print material, infill, and mounted hardware (driver, encoder, MCU, gyroscope/accelerometer, wiring), gives m_b, m_w, I_b, I_w directly, while l and l_b follow from the plate's design geometry. The damping coefficients C_b and C_w cannot be obtained from CAD, since they depend on bearing friction, wiring drag, and motor cogging; nominal values are used for the initial control design and replaced by the measured values below once the prototype is fabricated.

Table 7: Parameter determination method (symbols per Table 6).

Symbol	Determination method
l	Calculated during design
l_b	Calculated during design
m_b	CAD mass properties
m_w	CAD mass properties
I_b	CAD mass properties
I_w	CAD mass properties
C_b	Estimated (measured on the 1-DoF rig, below)
C_w	Estimated (measured on the 1-DoF rig, below)
K_m	Motor datasheet

1-DoF plant identification. Four parameters are measured on the single-axis rig once it is built, each by a method chosen so that the quantity of interest dominates the measurement rather than appearing as a small correction to something else. These measured values replace the CAD/datasheet estimates of Table 7 before the control gains used on hardware are finalised (Section 7).

Tolerances are stated with the measured values, since a parameter without a tolerance cannot support the robustness sweep of Section 7 that consumes it. Loop latency is measured through the complete path rather than estimated from the 1 kHz loop rate, because the scheduling and bus-transport contributions are the ones a rate figure omits.

3.4 Generalization to the 3D Cube

The dynamics, identification workflow, and verification approach validated on the 2D model of the first prototype extend to the second and final prototype: the full three-dimensional cube, where three orthogonal reaction wheels must stabilise the body about its edge and corner equilibria.

3.4.1 Full-Body Equations of Motion

With $R \in SO(3)$ the cube's orientation and $\omega_b \in \mathbb{R}^3$ its body-frame angular velocity, the attitude kinematics are

$$\dot{R} = R\,[\omega_b]_\times, \tag{3}$$

where $[\cdot]_\times$ is the skew-symmetric cross-product matrix. The assembly's angular momentum about the pivot, expressed in the body frame, generalizes the single-axis momentum term of the 2D model:

$$H = I\omega_b + \sum_{i=1}^{3} I_{w,i}\,\omega_{w,i}\,e_i. \tag{4}$$

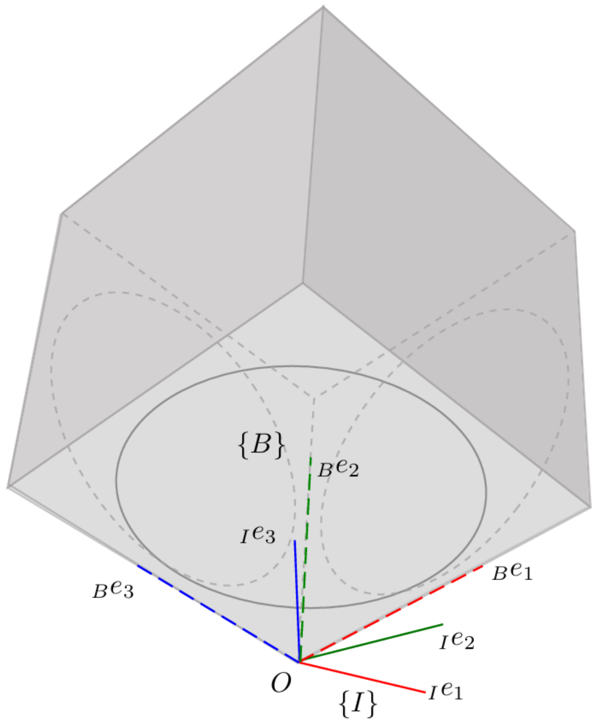

Figure 7: Body frame $\{B\}$ and inertial frame $\{I\}$ of the three-axis cube, both origined at the corner contact point O. The body axes $_Be_1, _Be_2, _Be_3$ coincide with the wheel spin axes e_i of Table 8, each directed outward from O along a cube edge; $\{I\}$ is fixed in space and coincides with $\{B\}$ at the corner-up equilibrium, so $R = \mathbb{1}$ there.

Euler's equation gives

$$\dot{H} + \omega_b \times H = \tau_g - C_b\,\omega_b, \tag{5}$$

and each wheel obeys

$$I_{w,i}\big(\dot{\omega}_{b,i} + \dot{\omega}_{w,i}\big) = T_{m,i} - C_{w,i}\,\omega_{w,i}, \qquad i = 1, 2, 3. \tag{6}$$

Table 8 defines all symbols.

The spin axis e_i is directed outward from the corner contact point O along the corresponding cube edge (Figure 7). Each wheel's motor/encoder positive-rotation direction is fixed by its mounting inside the inner structure and runs opposite to e_i, so the raw encoder rate and the commanded driver torque must be sign-flipped before they enter the equations above:

$$\omega_{w,i} = -\,\omega_{w,i}^{\text{enc}}, \qquad T_{m,i} = -\,K_{m,i}\,u_i^{\text{cmd}}, \tag{7}$$

so that H, Euler's equation and the per-wheel equation are all expressed consistently in the e_i-positive convention rather than the driver's native one.

Restricting motion to a single vertical plane with two wheel torques set to zero recovers the 2D model exactly, with $I_b, l_b, m_b, m_w, I_w, C_b, C_w$ the corresponding single-axis projections of $I, r_{cm}, m, I_{w,i}, C_{w,i}$.

3.4.2 Edge vs. Corner Equilibria & Linearisation

Jump-up proceeds in two stages: Flat-to-Edge, where one wheel brakes, followed by Edge-to-Corner, where a second wheel brakes. **Edge balancing** constrains the body to a single axis, so the dynamics reduce to the 2D form of Section 3.2.1, with m_b, l_b, I_b replaced by their edge equivalents.

Table 8: Full-body (3D) model variables.

Symbol	Description
R	Cube orientation relative to world frame
ω_b	Body angular velocity, body frame ($\in \mathbb{R}^3$)
$\omega_{b,i} = e_i^{\mathrm{T}} \omega_b$	Body angular velocity about wheel axis i
$\omega_w = (\omega_{w,1}, \omega_{w,2}, \omega_{w,3})$	Wheel angular velocities relative to body
e_i	Spin axis of wheel i (face normal)
I	3×3 inertia tensor, body+wheels, about pivot
H	Total angular momentum about pivot
$\tau_g = m\, r_{cm} \times (R^{\mathrm{T}} \hat{g})$	Gravity torque about pivot
m	Total assembly mass
r_{cm}	Body-frame vector, pivot to assembly CoM
$\hat{g}$	World vertical unit vector
C_b	3×3 body damping matrix
$I_{w,i}, C_{w,i}$	Inertia, damping of wheel i
$T_{m,i} = K_{m,i} u_i$	Motor torque, wheel i
$u_i, K_{m,i}$	Motor current, torque constant, wheel i

Corner balancing is genuinely 3D: all three wheels contribute across three rotational DOF. Linearising about the corner-up equilibrium $(\theta, \omega_b, \omega_w) = (0, 0, 0)$, with $\theta \in \mathbb{R}^3$ the small orientation error, gives the same state-space form as the 2D case:

$$\dot{x}(t) = Ax(t) + Bu(t), \qquad x := (\theta, \omega_b, \omega_w) \in \mathbb{R}^9, \tag{8}$$

with $A \in \mathbb{R}^{9\times 9}$ and $B \in \mathbb{R}^{9\times 3}$ built from I, r_{cm}, and the per-wheel parameters $I_{w,i}, C_{w,i}, K_{m,i}$ (Table 8). This linear model, together with its 2D counterpart of Section 3.2.1, is the plant the balancing controller of Section 7 is designed on.

4 Preliminary Sizing

This section confirms, before fabrication, that a cube of a candidate mass and edge length can perform the jump-up manoeuvre while carrying its own components. It covers: the sizing requirements (Section 1.3); the physics converting a target manoeuvre into a required wheel inertia (Section 4.1); the packaging constraint linking cube size to wheel size (Section 4.2); and the wheel design meeting that target (Section 4.3).

Provenance of the numbers in this section. Every figure reported below — the wheel-inertia target, the ballast station count, the wheel mass and inertia, and the cube mass budget — is produced by a single preliminary sizing program, `analysis/SIZING.py`, rather than assembled by hand from separate calculations. The program implements the jump-up dynamics in full (the ideal momentum transfer of Equation (10), the contact-case geometry factors of Table 9, the brake-amplification factor β and the resulting inertia target of Equation (11)), the wheel configuration (ring and spoke geometry, the radial M6 ballast stations with their parallel-axis and own-axis inertia terms, the axial-fit and edge-gap feasibility rules), and the cube mass budget itemised from the bill of materials with the structural allowance carried explicitly. Because the mass budget feeds back into the inertia target, the program is run to convergence rather than evaluated once, and the values quoted here are those of the converged pass. They are therefore reproducible: re-running the program with the inputs of Table 10 regenerates every number in this section, and changing an input regenerates all of them consistently rather than leaving the tables to be reconciled by hand. The figures remain *preliminary* in the sense that the structural line is still an allowance pending CAD closure, not in the sense of being approximate.

Sizing and control are solved together, not sequentially. Wheel inertia sets the angular momentum and torque available to the controller for arresting a fall — a wheel sized without the control law in mind can be mechanically sufficient yet dynamically useless. Conversely, the control law's saturation behaviour (how aggressively it must desaturate stored momentum, how much recovery angle it must cover) sets how much sizing margin is needed. Each depends on the other, so Section 4.2's closure procedure is a fixed-point iteration between them, not a one-pass calculation handed off afterward.

4.1 Jump-Up Feasibility Analysis

The analysis runs in one direction, which is worth stating before any equation appears. Two quantities enter as *inputs*: the cube mass M and edge length L, fixed by the packaging closure of Section 4.2; and the maximum wheel speed $\omega_{\max}$, a design choice bounded by what the drive can sustain in service (*Selection of* $\omega_{\max}$, below). One quantity leaves as the *output*: the per-wheel inertia $I_{w,\text{target}}$ the manoeuvre demands, which the wheel of Section 4.3 must then deliver through its diameter and ballast. Wheel speed is never solved for during sizing — it is set, and the wheel is sized around it.

Origin of the requirement. The physics comes from the jump-up collision itself, modelled as a perfectly inelastic transfer between a spinning wheel and the body. Conservation of angular momentum at impact and of energy from impact to edge-standing give a closed form for the speed a *given* wheel would need:

$$\omega_w^2 = (2-\sqrt{2})\,\frac{I_w + I_b + m_w l^2}{I_w^2}\,(m_b l_b + m_w l)\,g. \tag{9}$$

Equation (9) is the inverse of the sizing question rather than the sizing path: it presumes a wheel whose inertia I_w is already known, which at this stage is precisely the unknown. It serves two other purposes. Before fabrication it is a sanity check — a candidate wheel can be substituted to confirm that the speed it implies lies inside the motor's operating range at the available bus voltage, without excessively reducing the balancing recovery angle. After fabrication it becomes a verification relation, evaluated on the CAD- and test-derived $I_b, I_w, m_b, l_b, m_w, l$ of the built prototype (Section 3.3), with ω_w refined experimentally to absorb deviations from the ideal-collision assumption. The ω_w of Equation (9) and the $\omega_{\max}$ of Section 4.1.1 are the same physical wheel speed on opposite sides of one relation: here it is solved for; there it is fixed so that inertia can be solved for instead.

The braking mechanism uses a servo-driven barrier that strikes a bolt head on the wheel to stop it abruptly. Servo response time bounds the validity of the "instantaneous stop" assumption; a slower stop yields a lower effective post-impact body velocity than predicted, compensated through the same tuning process.

4.1.1 Torque-Limited Wheel-Inertia Target

Sizing follows the forward chain of Figure 8: cube parameters fix an ideal wheel momentum, a geometric/contact factor inflates it to a design momentum, and the design momentum converts to the target wheel inertia that Section 4.3 sizes the wheel geometry against, with $\omega_{\max}$ held fixed at every step.

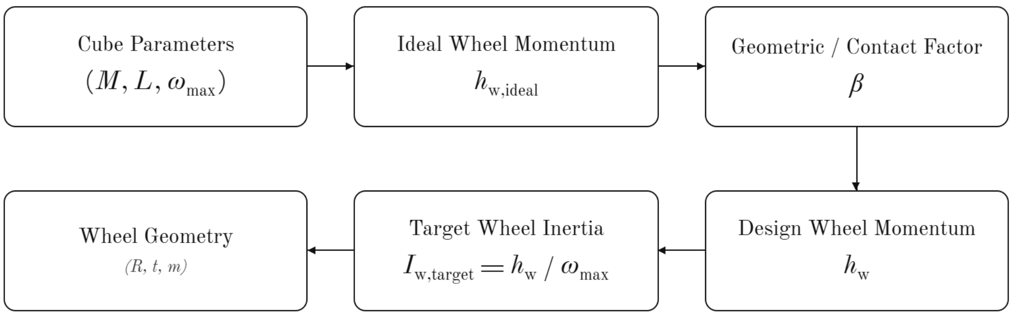

Figure 8: Wheel-inertia sizing chain.

Equation (9) assumes an instantaneous, lossless momentum transfer; the momentum budget below relaxes that assumption by accounting for finite brake torque. Its first step is an ideal per-wheel angular momentum:

$$h_{w,\text{ideal}} = \sqrt{\eta}\,\frac{\sqrt{2g\,\lambda\,\kappa}}{n}\,ML^{3/2}, \tag{10}$$

where M, L, g are cube mass, edge length, and gravity, η is an energy margin, and κ, λ, n are dimensionless geometric factors set by the contact case (Table 9). The corner case is more demanding and governs the design.

Table 9: Contact-case geometric factors.

Case	κ (pivot inertia)	λ (CoM rise)	n (momentum transfer)
Edge	$2/3$	$1/\sqrt{2} - 1/2$	1
Corner (governs design)	$11/12$	$\sqrt{3}/2 - 1/2$	$\sqrt{2}$

A finite brake torque τ_b cannot deliver the impulse instantaneously and must exceed the gravitational tip-over torque $\tau_g = MgL/2$. This loss is captured by a brake-amplification factor

$$\beta = \frac{\tau_b}{\tau_b - \tau_g}, \qquad h_w = \beta\, h_{w,\text{ideal}}, \qquad I_{w,\text{target}} = \frac{h_w}{\omega_{\max}}, \tag{11}$$

where $\omega_{\max}$ is the maximum practical wheel speed. Equations (10)–(11) are the sizing relations used at step 4 of Section 4.2's closure procedure: given a candidate mass and edge length, they return the inertia the wheel of Section 4.3 must meet. Their inputs, recorded in Table 10, are not fixed in advance — M and L are outputs of the packaging closure, not assumptions preceding it.

Table 10: Torque-limited wheel-inertia target (corner case). Inputs are fixed by the closure procedure of Section 4.2; the table is populated on each iteration. Values are those of the converged iteration.

Symbol	Description	Value
M	Cube mass	1.6638 kg
L	Cube edge length	149 mm
η	Energy margin	1.05
τ_b	Brake torque per wheel	5 N m
$\tau_g = MgL/2$	Gravitational tip-over floor	1.216 N m
β	Brake-amplification factor	1.321
$\omega_{\max}$	Max wheel speed (design)	5500 rpm*
h_w	Required angular momentum	$0.2351\,\mathrm{kg\,m^2/s}$
$I_{w,\text{target}}$	**Target wheel inertia**	$4.0813 \times 10^{-4}\,\mathrm{kg\,m^2}$

* Lowered from the 6000 rpm of the previous iteration. Since $I_{w,\text{target}} = h_w/\omega_{\max}$, a lower speed *raises* the inertia requirement and so demands more ballast — here six stations per wheel rather than three, costing 63.8 g of cube mass. What is bought in exchange is 500 rpm of relief on the drive and, because the station count rounds to multiples of three, a design point that sits on a wider tolerance band: the six-station wheel holds across 200.8 g of structure allowance (Section 4.2) where the three-station wheel held across roughly 46 g. No margin should be quoted at any other speed without re-running the closure, since the station count changes in steps rather than continuously.

Two features of these relations govern how the iteration behaves:

- $h_w \propto ML^{3/2}$ — growing the cube to fit a larger wheel also raises the requirement; convergence needs achievable inertia ($\sim D_w^2$) to grow faster than this.
- β diverges as $\tau_b \to \tau_g$ — a brake only marginally above the tip-over torque cannot deliver the required impulse at any wheel speed. τ_b must therefore be checked against $\tau_g = MgL/2$ every iteration, keeping a comfortable ratio τ_b/τ_g rather than a bare positive margin.

Selection of $\omega_{\max}$. Maximum wheel speed is a design choice, not a datasheet figure: it enters sizing directly through $I_{w,\text{target}} = h_w/\omega_{\max}$, so a higher assumed speed permits a proportionally smaller wheel. It is the primary relief variable when the envelope caps the wheel below the requirement (step 5, Table 11) — but it must reflect what the drive can sustain in service, not the no-load figure $K_v V_{\text{bus}}$. Three effects separate the two:

- winding resistance, which costs speed under the torque needed to spin the wheel up;
- reduced phase voltage available under sinusoidal field-oriented modulation;
- pack voltage sag through the discharge curve, taken at its lower end so the manoeuvre stays available on a partly-depleted pack, not only a fresh one.

The design value sits below the resulting ceiling, leaving margin for motor tolerance, bearing and aerodynamic drag, and residual uncertainty in M. Back-EMF should remain a comfortable fraction of usable bus voltage, and the electrical frequency $\omega_{\max} p/2\pi$ (at p pole pairs) must sit inside the driver's commutation capability — both to be confirmed once motor and pack are selected.

Selection of τ_b. Unlike the other inputs, τ_b cannot be read from a datasheet, because the brake does not generate it directly. The mechanism is a servo-driven barrier struck by a bolt head on the wheel: the servo only *holds* the barrier in the wheel's path, while the wheel's own momentum does the work of the collision. The servo's stall torque is therefore not τ_b and need not approach it — the TowerPro MG92B (Table 26) is rated near 0.29 N m, about one seventeenth of the sizing value. The servo only resists the barrier being deflected; it does not decelerate the wheel.

Since the stop is a collision rather than a friction clutch, τ_b is fixed entirely by how quickly angular momentum is destroyed,

$$\tau_b = \frac{h_w}{\Delta t}, \tag{12}$$

so quoting a brake torque is equivalent to asserting a stopping time: at the present operating point, $\tau_b = 5\,\mathrm{N\,m}$ implies $\Delta t \approx 47\,\mathrm{ms}$, whereas a bolt head striking a printed barrier should arrest the wheel in a few milliseconds. The design value is therefore expected to be conservative, which is the intended direction — since $\beta = \tau_b/(\tau_b - \tau_g)$ falls toward unity as τ_b grows, underestimating τ_b inflates the inertia requirement while overestimating it leaves the wheel undersized.

One consequence carries into the structural design: τ_b is also a *load*, and the sizing value should not be reused to size the barrier, bolt, servo bracket, or wheel spokes — a genuine few-millisecond stop puts an order of magnitude more torque through that load path. Until measurement settles the question, the structure is checked against a short-Δt case while the dynamics use $\tau_b = 5\,\mathrm{N\,m}$. The spin-down test (WBS task E10) closes this gap, reporting both the impulse-equivalent mean $\bar{\tau}_b = h_w/\Delta t$ (the sizing quantity) and the peak torque (the structural one); the two may differ several-fold, and the logging rate must resolve a few-millisecond event without smoothing the peak away.

4.2 Sizing Constraints and Closure Procedure

Cube edge length L and reaction-wheel inertia I_w are not independent choices — they are two ends of a single constraint chain, beginning with what must physically fit inside the cube. This section states that chain and the order in which it closes. No dimensions are committed here: component selection remains open (Sections 5, 6), and the numbers follow from the procedure rather than preceding it.

The chain runs in the order of Table 11. Internal volume is dominated by the LiPo pack and the subframe carrying it (locating the three motors on mutually perpendicular axes through a common centre, together with drivers, flight controller, IMU, and brake servos). The pack is incompressible and sits near the geometric centre to keep the centre of mass on the body diagonal, so L is the smallest edge length containing that volume plus wall thickness and contact features. L then caps wheel diameter D_w, and with it the achievable inertia $I_w \sim m_w D_w^2/4$ — ballast cannot rescue a wheel short of its maximum diameter, since it necessarily sits inboard of the ring it displaces.

Against that stands the jump-up requirement (Section 4.1.1),

$$I_{w,\mathrm{target}} = \frac{h_w(M, L)}{\omega_{\max}}, \qquad h_w \propto M L^{3/2},$$

which rises with both M and L. Two consequences close the loop: enlarging the cube to gain wheel diameter also raises the requirement, but achievable inertia grows faster than the requirement (D_w^2 against $L^{3/2}$), so a short wheel at a given L can be closed by a modest increase in L; and any mass added this way feeds back through $h_w \propto M$ and through the tip-over torque $\tau_g = MgL/2$ the brake must exceed — which is why the procedure terminates at a fixed point rather than in one pass.

Table 11: Sizing closure procedure. Steps are executed in order; a failure at step 4 returns to step 3 or step 1 as indicated.

Step	Determines	Governing consideration
1	Battery pack geometry	Capacity and discharge rating for the jump-up duty cycle
2	Inner clear volume V_{min}, hence minimum L	Pack + motor subframe + driver/controller stack, non-intersecting; pack near centre
3	Maximum wheel diameter D_w, hence achievable I_w	Clearance to subframe, motor body and the two orthogonal modules
4	Requirement $I_{w,\mathrm{target}} = h_w(M, L)/\omega_{\mathrm{max}}$	Corner jump-up, Eqs. (10)–(11)
5	Accept or iterate	If $I_w < I_{w,\mathrm{target}}$: raise L (step 2), raise ω_{max}, or trim M; re-enter at step 3
6	Mass reconciliation	Updated M feeds back into step 4; iterate to a fixed point

This procedure is implemented formally in `analysis/SIZING.py` and run to convergence by hand.

The two contact cases are not equally demanding: the corner equilibrium needs a larger centre-of-mass rise and a less favourable momentum-transfer geometry than the edge equilibrium, so at a common ω_{max} it demands the greater wheel inertia (Table 9). Step 4 is therefore always evaluated for the corner case — a wheel meeting the corner requirement meets the edge requirement with margin. Table 12 records both once the procedure closes.

Table 12: Required per-wheel inertia by contact case, evaluated at $\omega_{\mathrm{max}} = 5500\,\mathrm{rpm}$ for the converged $M = 1.6638\,\mathrm{kg}$. Only the geometry factor differs between the cases: τ_g and β are properties of the cube and the brake, not of the feature it stands on.

Contact case	h_w ($\mathrm{kg\,m^2/s}$)	$I_{w,\mathrm{target}}$ ($\mathrm{kg\,m^2}$)	Relative
Edge	0.2133	3.7026×10^{-4}	90.7 %
Corner (critical)	0.2351	4.0813×10^{-4}	100 %

The corner case demands 9.3 % more inertia than the edge case, so sizing on it covers the edge case with margin. The selected wheel delivers $4.3356 \times 10^{-4}\,\mathrm{kg\,m^2}$ (Table 14) — 106.2 % of the corner requirement and 117.1 % of the edge requirement.

Table 13 is the bookkeeping side of steps 5 and 6: populated as components are selected and re-totalled each iteration, since the total is both an input to $I_{w,\mathrm{target}}$ and an output of the design.

4.3 Reaction-Wheel Sizing

The wheel must supply the inertia $I_{w,\mathrm{target}}$ demanded by the corner jump-up (Section 4.1.1), fit within the envelope's outer diameter (step 3, Table 11), and stay light enough not to dominate cube mass, since its own mass feeds back into $I_{w,\mathrm{target}}$ through M. These three demands pull against each other, so the wheel architecture is chosen to be trimmable after fabrication rather than needing to be right the first time.

Architecture: ballasted ring. Rather than scaling a solid disc, the wheel is a printed *ring plus spokes* tuned by *ballast*: standard M6 steel bolt-and-nut sets fitted through round clearance holes drilled *radially* through the ring's cross-section, from the inner bore surface to the outer

Table 13: Mass budget (per cube), rolled up from the bill of materials by `analysis/SIZING.py` (Stage 6). The total is an input to the inertia requirement, not only a result: the value below is the *converged* one, which the closure reproduces exactly ($\delta = 0.0\,\mathrm{g}$) against the M assumed in Table 10.

Subsystem	Mass (g)	Share	Basis
Reaction wheels ($3 \times 170.37\,\mathrm{g}$)	511.1	30.7 %	Sized (Sec. 4.3)
Motors and brake servos (3× each)	245.4	14.7 %	Datasheet
Power (battery, XT60, regulator)	279.0	16.8 %	Drives envelope (step 1)
Electronics (drivers, MCU, encoders, IMU)	95.2	5.7 %	Datasheet, as-weighed
Wiring, perfboard and passives	45.5	2.7 %	Estimated
Subtotal — specified hardware	**1176.2**	70.7 %	Measured or datasheet
Structure (subframe, frame, fasteners)	487.6	29.3 %	**Allowance, pending CAD**
Total M	**1663.8**		**Converged fixed point**

The structure line is an *allowance*, not a measurement, and at roughly 29 % of the total it remains the dominant uncertainty in the budget. It has been revised *upward* from the 374 g of the previous iteration to 487.6 g (subframe 199.5 g, frame 239.9 g, fasteners/harness 48.2 g, held at the same relative split), the earlier figure having proved optimistic against the frame and subframe now in CAD. The electronics line rises over the same revision, 74.2 g to 95.2 g, on as-weighed rather than bare-board figures — the flight computer carries headers, each encoder breakout is quoted with its steel-backed ring magnet, and the telemetry radio with its antenna and pigtail. The allowance is reported separately from the specified hardware for that reason: the subtotal is what the design currently knows, and the total is what it currently projects. Closing the CAD model replaces the allowance with a real figure and re-enters the loop at step 6. The present design point tolerates that: the six-station wheel holds anywhere between 364.1 g and 564.9 g of structure allowance, so the budget can absorb an overrun of about 77 g before the station count steps again.

rim surface — a true ID-to-OD through-hole, not one drilled through the wheel's face. Each hole's depth is fixed at the ring's radial width, and its centroid sits at the mean radius $R_{\mathrm{mean}} = (D_w + D_i)/4$ by construction, with no free pitch-circle radius to choose. Hole diameter is instead capped by the wheel's axial thickness, with a printed wall retained on both faces. Each station removes plastic and adds steel at very nearly the largest available radius, so the net inertia gain per station is large, and the target is reached simply by choosing how many stations to populate.

The bolt is also the retention feature: the nut threads on and clamps the ring, so nothing depends on a plastic pocket floor. Each station carries the centrifugal load $F = m_{\mathrm{hw}}\,\omega_{\max}^2\,R_{\mathrm{mean}}$ of its own hardware, reacted by the bolt in tension and by bearing of the head and nut faces — now against the ring's curved ID and OD surfaces rather than a flat face, so both are spot-faced flat (or fitted with a curved washer) to seat squarely. A nyloc or threadlocked nut is specified so vibration cannot back it off.

The geometry is defined by the quantities of Table 14, all of which follow from the envelope closure of Section 4.2 and are recorded there once fixed.

The ballast is characterised *net* of the plastic each hole displaces: m_b and ΔI_b are the increments the installed bolt-and-nut adds over the radial cylinder of material it replaces, evaluated at the forced centroid radius R_{mean}. Both corrections use the parallel-axis theorem about the spin axis, retaining each feature's own-axis term rather than treating it as a point mass — here the *transverse* solid-cylinder term $r_h^2/4 + w_r^2/12$, since the hole's axis is radial (perpendicular to the spin axis), unlike a face-drilled hole. Sizing then reduces to solving

$$I_w(D_w, N) = I_{w,\mathrm{ring}}(D_w) + N\,\Delta I_b \geq I_{w,\mathrm{target}}$$

for the smallest admissible N at each candidate D_w, subject to the fit and mass checks below.

Sizing calculator and iteration procedure. These relations are implemented as a parametric calculator mapping the geometry of Table 14 onto the mass and inertia of one wheel: it

Table 14: Reaction-wheel geometry and ballast parameters. Values are the current design point of the iteration described below; they are provisional while the structural mass of Table 13 remains an allowance.

Symbol	Parameter	Value
D_w	Ring outer diameter (fixed by envelope)	120 mm
D_i	Ring inner diameter ($D_w - 2w_r$)	100 mm
w_r	Ring radial width (= ballast hole depth)	10 mm
t	Wheel axial thickness (caps hole diameter)	20 mm
n_s, w_s	Spoke count and width	3, 10 mm
ρ_p	Printed-material density (PET-CF)	1290 kg/m^3
d_h	Ballast hole diameter (M6 clearance)	6.4 mm
R_{mean}	Ballast centroid radius (forced, $= (D_w+D_i)/4$)	55.0 mm
m_{hw}	Hardware mass per station (bolt + nut)	7.5 g
m_b	*Net* mass added per station	7.085 g
ΔI_b	*Net* inertia added per station	2.151×10^{-5} kg m^2
N	Station count reaching $I_{w,\text{target}}$ (of $N_{\max} = 37$)	6
m_w	Resulting wheel mass	170.37 g
I_w	Resulting wheel inertia	4.3356×10^{-4} kg m^2
	as a fraction of $I_{w,\text{target}}$	106.2 %

sweeps candidate outer diameters and, for each, returns the required station count, the count the gap rule admits, resulting mass and inertia, and the pass/fail state of every constraint. It is not an optimiser, and the reported design point was not found by one.

Outer diameter D_w is not free to choose — it is imposed by the packaging envelope, since three orthogonal wheel modules and the battery must fit within L — so the sweep is retained for what it reveals, not what it selects. Left free to minimise wheel mass, the calculator prefers a larger ring, since $I_{zz} \sim mR^2$ makes mass at large radius buy inertia more cheaply: the unconstrained optimum at the present target is 140 mm (151.81 g/wheel, no ballast needed), so the envelope constraint at 120 mm costs about 18.6 g per wheel (55.7 g over three). The penalty has roughly doubled since the previous iteration, because the higher inertia target now has to be met with steel at R_{mean} rather than with ring diameter. This is the price of a wheel that fits, and a wheel that does not fit is not a candidate.

The procedure followed is the fixed-point iteration of Table 11, run by hand: a cube mass is estimated and a mass budget assembled; the jump-up analysis converts that mass into an inertia requirement; the wheel is sized against it subject to volume and hole-pattern constraints; and the resulting wheel mass feeds back into the budget that started the loop. The cycle repeats until the geometry simultaneously fits the internal volume, meets the inertia requirement, and reproduces the cube mass it assumed — convergence is what closes the design, with step 6 of Table 11 returning approximately zero.

Selection among admissible diameters is made on minimum *wheel mass*, not minimum diameter.

5 CAD Design

This design is preliminary: final actuator and sensor specifications are not yet confirmed, so the mechanical architecture is kept modular and largely parametric (Section 1.1). This section states the mechanical concept and the CAD conventions the build follows, for both the first prototype (the single-face 1-DoF rig) and the second and final prototype (the three-axis cube), and the checks that confirm the model and the fabricated parts agree. It is deliberately kept to the level of decisions and interfaces: the part-by-part model tree, the manufacturing drawings, the fabrication settings and the assembly sequence are recorded in Appendix B, so that a change to a bracket does not require an edit to the design rationale, and the rationale is not buried in geometry.

5.1 Mechanical Design and CAD Construction

5.1.1 CAD Methodology and Model Conventions

The model is built top-down from the envelope closure of Section 4.2 rather than bottom-up from individual parts, so that the edge length L and the wheel diameter D_w propagate to every dependent feature when an iteration changes them. Three conventions make that propagation reliable, and they are stated here because they constrain how every part in Appendix B is modelled:

1. **A single skeleton drives the assembly.** L, D_w, the motor axis offsets and the wall thickness live in one parameter table referenced by every part; no downstream part carries a hard-coded dimension that duplicates one of them.
2. **The body frame is the CAD origin.** The assembly origin is the cube's geometric centre, with axes on the three motor spin axes, so that CAD-reported mass properties are directly the I and r_{cm} of Table 8 without a change of basis.
3. **Mass properties are an output of the model, not an annotation.** Every part carries its assigned material and print density, so the assembly mass roll-up is the quantity fed back into Table 13 at step 6 of Table 11.

Convention 3 makes the iteration auditable: the mass budget and the inertia requirement are read from the same model, so a design that closes analytically cannot fail to close in the CAD without the discrepancy being detected.

5.1.2 1-DoF Prototype Build

The single-face rig is the first article built and the bench on which the parameters of Table 7 are measured. Its structure is a 3D-printed plate sized to one cube face, a central motor and wheel mount, a corner bearing defining the single rotational degree of freedom, and encoders and other hardware. Note that the servo (brake system) is not included in the 1-DoF prototype as its purpose is validating the control loop on a more well-understood problem. Detailed geometry and the drawing set are in Appendix B.3.

The rig as built is shown in Figure 9. The plate is mounted diagonally, so that the pivot bearing sits at the lowest corner and the wheel axis passes through the plate centre: this places the centre of mass directly above the contact point at the upright equilibrium, which is the configuration linearised in Section 7. The two printed uprights either side of the frame are hard stops, limiting travel to a few degrees past the point of no return so that a failed balancing attempt ends against a stop rather than on the bench.

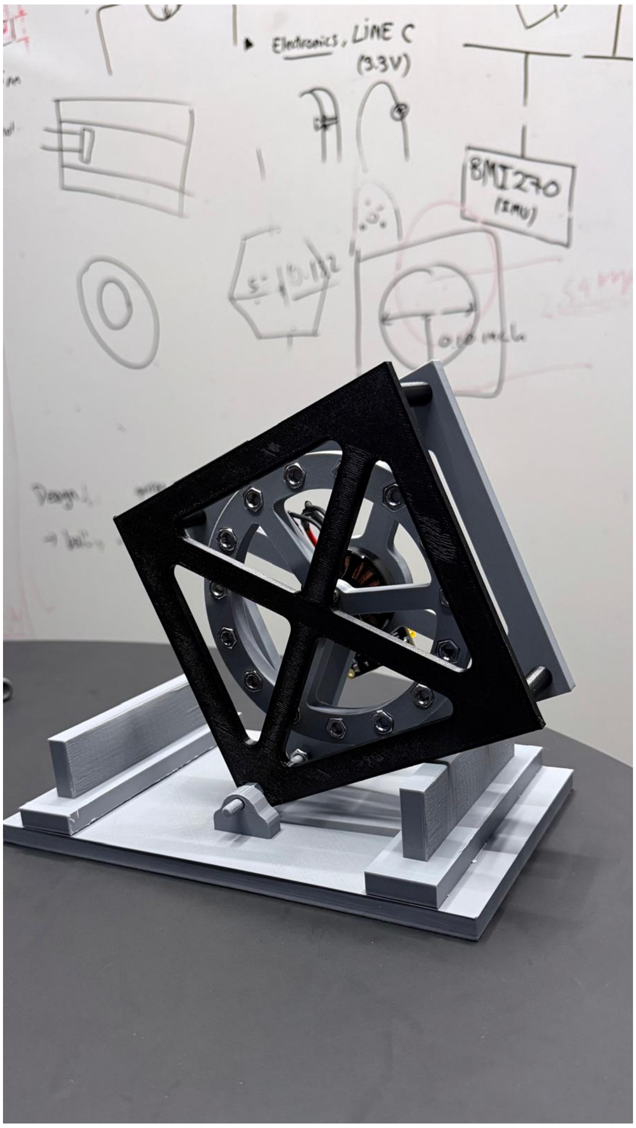

Figure 9: The single-face prototype rig as built: one reaction wheel driven by a brushless motor on a diagonally-mounted printed plate, pivoting on a corner bearing against a printed base with hard stops. This is the article on which the plant parameters of Table 7 are measured and the single-axis controller of Section 7 is validated before the three-axis build. The drawing set is given in Appendix B.3.

5.1.3 3D Cube Frame

The frame is conceived in two parts, reflecting the constraint chain of Section 4.2. An *inner structure* locates the three motors on mutually perpendicular axes through a common centre and carries the battery pack, the three drivers, the flight controller, inertial measurement unit, and the brake servos; it is the element that fixes the minimum internal clear volume and therefore the edge length L. An *outer frame* carries the edge and corner contact features, protects the wheels, and reacts the motor mounting loads from the inner structure into the ground contact.

A first physical prototype of this frame has been printed and is shown in Figure 10. It is explicitly a *concept prototype*, not the flight article: it demonstrates the two-part architecture, the X-braced face pattern and the packaging of three orthogonal wheel modules within the cube envelope, and it serves as a physical check on clearances that are awkward to judge on screen. The final design is not frozen at the time this photograph was taken. Detailed geometry follows the envelope closure of Section 4.2, and the wheel dimensions in particular remain open pending the sizing iteration of Section 4.3 and confirmation of the final motor; the mass properties of the printed prototype are therefore not the mass properties assumed in the budget. The frozen general-assembly drawing is given in Figure 17 of Appendix B.4.

The split is also the main mechanical interface in the machine, and it is managed as one: the inner structure is designed against a fixed bolt pattern and keep-out envelope published to the outer frame, so the two can be revised independently as long as that interface holds. The controlled interfaces are listed in Table 15; the parts realising them, with their drawings, are in Appendix B.4.

Table 15: Controlled mechanical interfaces. Each is frozen once agreed, so that the parts on either side can be revised independently.

Interface	Between	Controlled quantities
Motor mount	Motor ↔ inner structure	Bolt circle, shaft height, concentricity to the wheel axis
Wheel hub	Wheel ↔ motor rotor	Bore, bolt pattern, axial location, balance datum
Frame joint	Inner structure ↔ outer frame	Bolt pattern, keep-out envelope, load path to the contact features
Electronics mount	Board ↔ inner structure	Board outline, mounting-hole pattern, keep-out heights, connector exit directions (Sec. 6.6)
Brake mount	Servo ↔ inner structure	Servo location, barrier throw
Battery retention	Pack ↔ inner structure	Pack envelope, retention method, centring tolerance to the cube centre
Contact features	Outer frame ↔ ground	Corner and edge geometry, material, replaceability

5.2 CAD and Dimensional Verification

The checks below split into those that run against the model, before anything is fabricated, and those that run against the physical part afterwards. The driving dimensions themselves are outputs of the packaging closure of Section 4.2 and are not verified here; what is verified is that the model realises them consistently and that the fabricated part matches the model.

Model checks (pre-fabrication). Interference and clearance verification on the wheel-to-subframe and inter-module clearances c_w and c_o of Table 27, through a full rotation of each

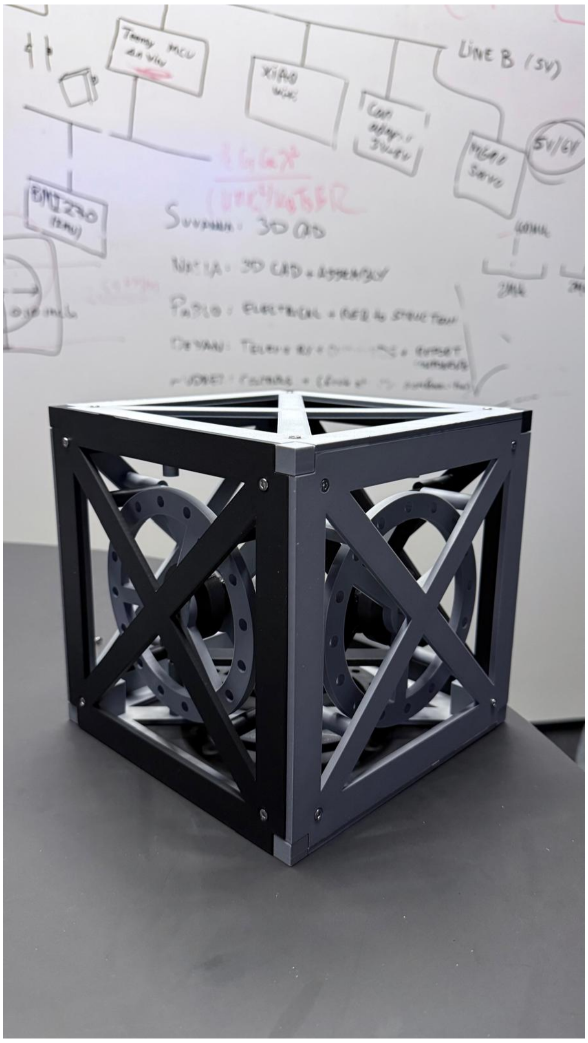

Figure 10: Printed concept prototype of the three-axis cube frame, showing the X-braced outer faces and the three orthogonal reaction-wheel modules carried by the inner structure. This article validates the two-part architecture and the internal packaging only: it is *not* the final design. Wheel geometry and the frame edge length L remain subject to the sizing closure of Sections 4.2 and 4.3, so its mass properties do not represent the budgeted values. The frozen design's general-assembly drawing is given in Figure 17 of Appendix B.4.

wheel and across all three axes simultaneously; orthogonality of the three motor axes; position of the centre of mass relative to the geometric centre, since the modelling of Section 3 assumes a known offset; and the mass roll-up against Table 13, which is the input to the sizing verification and must be re-run whenever it changes.

Part checks (post-fabrication). Dimensional acceptance of the ballast holes — diameter, pitch circle radius and the radial wall left on each side — since these are structural features rather than cosmetic ones and a print that is dimensionally out of tolerance at a hole is a wheel that cannot be trusted at speed; verification that each bolt is torqued and its nyloc nut engaged, the through-bolt being the retention feature; measured wheel mass against the CAD prediction; and static balance of the assembled wheel, which is the first opportunity to detect a ballast placement error before it appears as vibration at speed.

The modelled-versus-measured comparison is recorded in Table 31 of Appendix B.6 and is not duplicated here. A discrepancy found at that step is carried back into $I_{w,\mathrm{target}}$ and absorbed by ballast trim where the authority allows; where it does not, the finding returns to the closure procedure rather than being accepted.

6 Electronics

This section is presented at the level of the decisions that shape the electrical architecture: what the components are, how the whole system loop fits together, how power is distributed, how the buses are routed and terminated, and how the assembly is made safe to power on. The sheet-level schematics, the connector and pin tables, the harness drawings and the perfboard layout are collected in Appendix C, for the same reason the CAD detail is separated from the mechanical rationale of Section 5 — a connector re-pin should not require an edit here. The acceptance criteria and checks run against this architecture during bring-up (CAN-bus termination, rail voltages, sensor and driver calibration) are tracked against the electrical bring-up stage of the validation strategy, not repeated here.

6.1 System Architecture

Figure 11 shows the closed control loop that ties the four subsystems together. The Teensy 4.1 runs estimation and control at 1 kHz and issues torque commands over CAN-FD to three moteus-n1 drivers, each closing a 30 kHz field-oriented current loop on its MN4006 motor. The resulting reaction torque acts on the cube body; the body's attitude is sensed by the BMI270 IMU and the rotor angle by the MA600 absolute encoder, closing the loop back to the controller. The two loops run at different rates by design: the fast current loop is delegated entirely to the drivers, so the flight controller need only meet the 1 ms outer-loop deadline of Section 7.

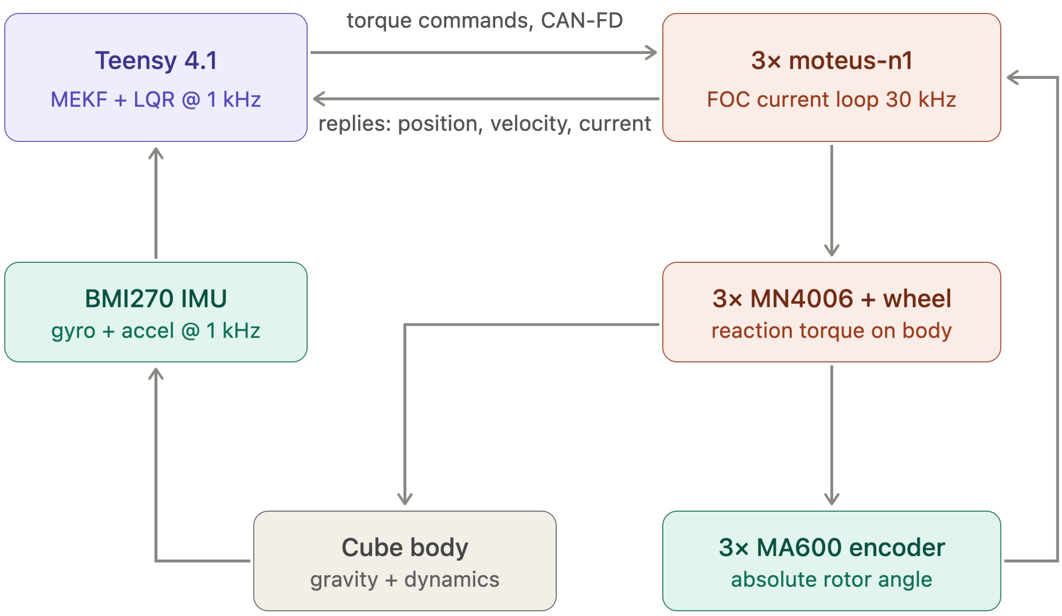

Figure 11: System architecture and control loop. Blue: processing; green: sensing; orange: actuation. Rates shown are the design targets of Section 7.

6.2 Components

The electronics and software are built around the Teensy 4.1 microcontroller, which communicates over the CAN-FD bus with the mjbots moteus-n1 driver that closes the current loop for the MN4006 brushless motor. Wheel speed is measured with the MA600 magnetic encoder on the motor shaft, and body tilt is estimated from the BMI270 IMU. Braking is actuated by a

TowerPro MG92B digital servo driven directly from the microcontroller, and the XIAO ESP32-C6 wireless module provides a telemetry and debugging link. Power is supplied by the 6S LiPo battery, stepped down through an LM2596 DC-DC converter for the logic-level electronics, with the motor drivers powered directly from the battery bus. Table 16 lists each component against the schematic block it belongs to (Appendix C.1); current ratings are given in Appendix C.2 rather than here, so that a part substitution does not require an edit to this section.

Table 16: Component overview: each part against the schematic block that carries it (Figure 18, Appendix C.1).

Component	Function	Schematic block
Teensy 4.1	Flight computer: estimator + control law	Flight computer and IMU
BMI270 IMU	Body attitude reference (tilt + rates)	Flight computer and IMU
mjbots moteus-n1 (3×)	FOC current/torque driver, inner loop	Power entry and protection; CAN-FD bus
T-Motor MN4006 (3×)	Reaction-wheel torque source, via driver	Power entry and protection
MA600 encoder (3×)	Wheel angle: commutation + state feedback	Encoder interfaces
TowerPro MG92B servo (3×)	Wheel brake actuator (jump-up)	Brake servo drive
XIAO ESP32-C6	Wi-Fi telemetry / command link	Telemetry link
LM2596 buck converter	5 V logic rail source	5 V logic rail
6S LiPo pack	Primary energy store	Power entry and protection

The full parts list is catalogued in Appendix A (Table 26). The supplied set is an electronics-only BOM; the reaction wheels, the frame and all printed parts are the team's design responsibility, and their fabrication is the leading schedule risk.

Table 17 states, block by block, the layout, mounting and connection constraints each part is checked against before bring-up: what governs its placement, and how the assembled machine confirms it. These are assembly and layout restrictions on the physical prototype, not schematic content, which is why they are stated here rather than repeated against each schematic block of Section 6.4.

Motor characterisation and operating envelope. Figure 13 shows measured torque and speed against phase current for the MN4006, taken from manufacturer bench data. The measurements were made with the motor driving propellers rather than reaction wheels, so the *load* is not representative of this application — an aerodynamic load rises with the square of speed, whereas a reaction wheel is an inertia with only bearing friction opposing it at constant speed. What does transfer is the motor's own electromechanical behaviour, since torque per ampere and back-EMF per unit speed are properties of the machine and not of what it is attached to. The curves are used here only in that sense: to bound the torque and current the actuator can be expected to deliver.

Three figures relevant to the design follow from the data. First, the torque–current relation in Figure 13(a) is close to linear ($r = 0.97$) with an incremental slope of $0.023\,\mathrm{N\,m/A}$ and a positive torque-axis intercept of about $0.08\,\mathrm{N\,m}$. The intercept is the expected signature of losses that do not reach the shaft — magnetising and iron losses plus bearing friction — so the useful torque

Table 17: Electrical layout, mounting and connection constraints, by block. Checked against the assembled prototype before and during bring-up.

Component / block	Layout, mounting & connection rule	Verification
Power entry and protection	Pack connector, anti-spark inrush path, E-stop and fuse in series, upstream of the split to the motor bus and logic rail	E-stop functional; fuse rated to worst-case bus current; anti-spark path intact
5 V logic rail	LM2596 buck converter off the motor bus, back-feed diode, bulk capacitance and per-device decoupling at each load	Converter rating vs. measured worst-case rail current (Appendix C.2); capacitor voltage ratings confirmed by position
CAN-FD bus	Linear chain across the three moteus drivers and the flight computer, short stubs, termination at the two physical ends only (Sec. 6.5)	Differential-pair scope capture; sustained error-frame count at full data rate with all nodes addressed
Encoder interfaces (3×)	Short SPI run from each driver to its own encoder; encoder positions fix the driver positions, not the reverse	Air gap ≤ 0.5 mm confirmed; read-back with no error flags, rechecked after mechanical assembly
Flight computer and IMU	Teensy 4.1 and BMI270 mounted at or as close as possible to the pivot; orientation set by packaging, corrected in software by the mounting rotation of Sec. 7.2	Static orientation check; noise floor confirmed with wheels spinning
Telemetry link	ESP32-C6 module fed off the logic rail, sized for its transmit-burst current	Transmit-burst return path verified off the logic rail
Brake servo drive	Servo PWM routed clear of the encoder SPI runs, one instance per axis	Stall-current headroom confirmed on the logic rail
Rail separation	Motor bus and logic rail kept separate downstream of the E-stop	Ripple measurement on the logic rail under motor load
Grounding	Single return topology; no motor return current through the sensing ground	Continuity and isolation check before first power-on

PLACEHOLDER — artwork pending

Assembled electronics of the Cubli: flight computer, driver stack, distribution board and harness in place inside the frame.

Figure 12: Assembled electronics of the Cubli. To be replaced with a photograph once the perfboard and harness of Section 6.6 are populated.

per ampere is the slope, and a constant inferred by simply dividing torque by current at a single point would overstate it (that ratio runs to 0.031 N m/A through the origin). Either way the motor produces at most 0.45 N m, an order of magnitude below the 5 N m brake torque assumed in Section 4.1.1. The two are not alternatives: the motor stores angular momentum in the wheel over about a second, and the brake destroys it in tens of milliseconds, so the impulsive torque comes from the collision and not from the drive. This is why braking is a separate mechanism rather than a motor function. Second, the highest measured point reaches 0.45 N m at 17.5 A, above the 16 A peak rating in Table 26, and that run terminated on the motor's thermal limit — so the top of this envelope is a short-duration capability, not a continuous one. Third, speeds cross the 5500 rpm design ceiling of Section 4.2 only under the lighter loads, which is consistent with that ceiling being set by bus voltage and derating rather than by any inability of the motor to spin faster.

6.3 Power Architecture

The power architecture is shown in Figure 14. A single 6S LiPo pack (22.2 V nominal, 1550 mA h) feeds everything through one XT60 connector, behind a combined E-stop and fuse, with an anti-spark path to limit inrush into the driver bulk capacitance. The XT60 is fixed by the pack rather than chosen: it is rated at $\approx$60 A, which covers the three-axis spin-up transient but with a smaller margin than a 90 A housing would give, so the transient is a measured acceptance item at first power-on rather than an assumed one. The bus then splits into two rails. The *motor bus* runs at pack voltage (22–25 V across the discharge curve) directly into the three moteus drivers, with bulk capacitance rated to at least 50 V to absorb regenerative braking transients—which are not incidental here, since the jump-up manoeuvre of Section 4.1 dumps the full wheel kinetic energy back into the bus. The *logic rail* is derived by an LM2596 buck converter stepping 22 V down to 5 V for the Teensy, the ESP32 telemetry link, the three brake servos and the IMU. Keeping the two rails separate downstream of the E-stop means motor-current ripple does not reach the sensing and processing electronics.

The rail-by-rail current budget that sizes the buck converter and the harness conductors is given in Appendix C.2 (Table 32), populated from datasheet ratings and completed from measurement during bench bring-up rather than from typical values alone, because the figure that matters for the 5 V rail is the coincident worst case — telemetry transmit burst, servo stall and full sensor

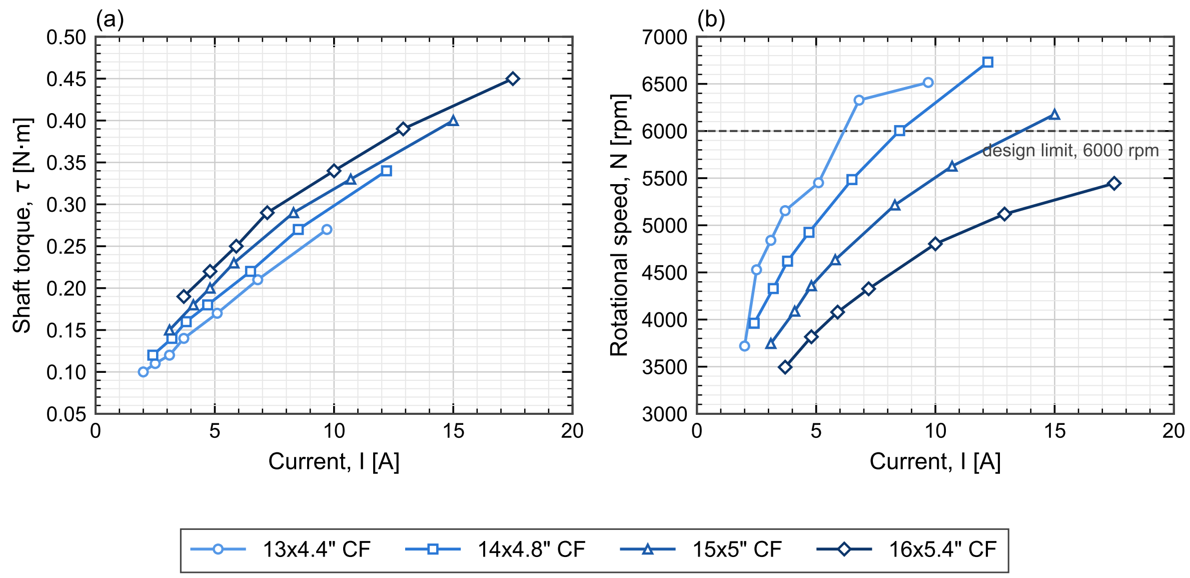

Figure 13: Measured MN4006 (a) shaft torque and (b) rotational speed against phase current, for one motor driving four different loads. Symbols are measured points at throttle settings of 50, 55, 60, 65, 75, 85 and 100 %; lines join successive points and are not fits. The dashed line in (b) is the 5500 rpm reaction-wheel design ceiling of Section 4.2. Loads shown are propellers, not reaction wheels: the curves characterise the *motor*, and the load-dependent spread between them does not carry over to this application. Manufacturer static bench data; the heaviest load reached the motor thermal limit at 100 % throttle.

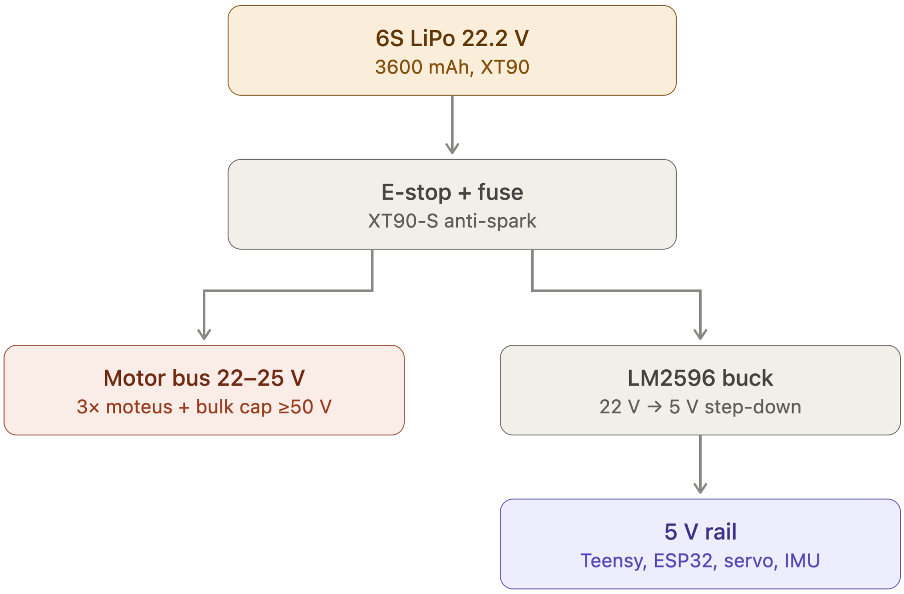

Figure 14: Power distribution tree. A single 6S pack feeds the motor bus directly and the 5 V logic rail through a buck converter, both behind a common E-stop and fuse.

load at once. The 5 V total is checked against the LM2596 rating, which derates without a heatsink.

6.4 Schematic and Interconnect Diagrams

The signal-level design is captured as one flat schematic (Figure 18, Appendix C.1), together with the connector and pin-assignment tables. The layout, mounting and connection constraints each block is checked against are given in Table 17 of Section 6.2, ahead of the schematic itself, and are not repeated here.

6.5 Bus Topology, Termination and Signal Integrity

The CAN-FD bus is the one electrical subsystem whose physical routing is a correctness requirement, and it is therefore fixed before the board outline is drawn. At the design data rate the bus must be a linear chain with short stubs and termination at exactly the two physical ends; a star topology or a mid-chain termination will pass a bench continuity check and then fail intermittently under wheel vibration, which is the hardest class of fault to diagnose once the cube is closed. The design rules and their verification are listed in Table 17 of Section 6.2, together with every other block's layout and connection constraints.

The encoder interfaces impose the complementary constraint. Each absolute encoder requires a short SPI run to its own driver and a tightly held air gap to its magnet, so the driver positions are dictated by the encoder positions, and those in turn by the motor axes. The physical layout order is therefore: motor axes fix the encoders, encoders fix the drivers, drivers fix the bus route, and only then is the board outline of Section 6.6 drawn around what remains.

6.6 Board Layout and Harness

The electronics are built on a double-sided perfboard rather than a fabricated PCB. A PCB would have a long lead time, and a perfboard allows for iteration and modification throughout the process. Its consequence is that layout discipline — rail separation, decoupling at each device, deliberate routing of the bus pair — must be enforced by hand and checked explicitly, since no design-rule check will do it.

The board outline, mounting-hole pattern, keep-out heights and connector exit directions constitute the electronics-mount interface of Table 15 and are frozen once delivered to the mechanical design: later electrical changes must fit inside that outline rather than move it. The harness — battery lead, driver drops, bus chain and servo drive — is fabricated to the drawings in Appendix C.4 and is labelled and strain-relieved before installation, because every connector in the machine sits on a structure that vibrates by design.

IMU mounting and axis alignment. The BMI270 is carried on Perfboard 1 alongside the Teensy (Figure 17), placed by wiring and packaging convenience rather than by axis alignment to the body frame $\{B\}$ of Figure 7. Its native sensor axes therefore do not coincide with $\{B\}$; every sample is corrected by the fixed mounting rotation of Section 7.2 before the estimator uses it, rather than by re-orienting the board to match the model.

6.7 Electrical Safety and Bring-Up Protocol

Two rules govern all electrical bring-up and are stated here because they are design constraints, not merely lab practice. First, every first-time power application is made on a current-limited bench supply, never on the battery; the battery is introduced only after the assembly has passed a full bench power-on. Second, the pack is never connected without the E-stop and fuse in circuit and the anti-spark path intact, since the driver bulk capacitance draws an inrush that will weld a bare connector.

The staged sequence — rails, then buses, then sensors, then open-loop motion, then closed-loop — and the acceptance criteria that close each step are tracked against the project's validation stage gates. The isolation checks that precede first power-on (rail-to-rail separation, bus-to-chassis, and confirmation that no low-voltage capacitor sits on the pack bus) are recorded in Appendix C as a checklist, because they are performed on every re-assembly and not only once.

7 Control Algorithms & Approach

This section derives the balancing controller from the linear models of Section 3, states the controller as executable logic rather than only as equations, describes the Simulink/Simscape environment used to exercise that logic against both prototypes' CAD geometry before it runs on hardware, and gives the onboard software architecture that carries it at 1 kHz in flight. Detailed pseudocode, the fixed gains and tuning parameters, the timing budget and the fault responses are collected in Appendix D.

7.1 LQR Baseline Control

Linearising about the upright equilibrium $(\theta_b, \dot{\theta}_b, \dot{\theta}_w) = (0, 0, 0)$ gives

$$\dot{x}(t) = Ax(t) + Bu(t), \qquad x := (\theta_b, \dot{\theta}_b, \dot{\theta}_w), \tag{13}$$

with A, B built from the parameters in Table 6. Zero-order-hold discretization at a rate set by the open-loop unstable pole, with a safety margin, gives $x[k+1] = A_d x[k] + B_d u[k]$. A discrete LQR gain $K_d = (K_{d1}, K_{d2}, K_{d3})$ minimises

$$J(u) = \sum_k x^{\mathrm{T}}[k]Qx[k] + u^{\mathrm{T}}[k]Ru[k], \tag{14}$$

giving $u[k] = -K_d(\hat{\theta}_b[k], \hat{\dot{\theta}}_b[k], \hat{\dot{\theta}}_w[k])$. The same design generalises directly to the corner-balancing 9-state model of Section 3.4: K_d becomes a 3×9 gain matrix, using the same discretization, cost function and offset filter below, with Q, R weighting all three tilt axes, body rates and wheel speeds.

The tilt estimate $\hat{\theta}_b$ is obtained from a gyroscope and accelerometer mounted at, or as close as mechanically feasible to, the pivot point rather than at the plate's centre. This placement is necessary because an accelerometer at radius r from the pivot measures gravity plus the tangential ($r\ddot{\theta}_b$) and centripetal ($r\dot{\theta}_b^2$) acceleration components induced by the plate's rotation, which corrupt the naive estimate $\tan^{-1}(-a_x/a_y)$ during fast motion (e.g., jump-up or aggressive recovery); placing the sensor at $r \approx 0$ makes both terms negligible. Tilt is estimated with a complementary filter combining gyro integration (drift-prone, unaffected by the terms above) with accelerometer inclination (drift-free, but unreliable during shocks):

$$\hat{\theta}_b[k] = (1-\gamma)\big(\hat{\theta}_b[k-1] + \dot{\theta}_{b,\text{gyro}}[k]\,\Delta t\big) + \gamma\,\theta_{b,\text{accel}}[k], \tag{15}$$

where $\theta_{b,\text{accel}}[k] = \tan^{-1}(-a_x[k]/a_y[k])$ is the instantaneous accelerometer-only estimate and $\gamma \in (0,1)$ is a small weight favoring the gyro term between corrections.

A constant tilt-sensor offset d, including residual bias from imperfect pivot placement or sensor misalignment, would otherwise appear as a nonzero steady-state wheel speed rather than a corrected angle. A low-pass filter state

$$x_f[k+1] = (1-\alpha)x_f[k] + \alpha\hat{\theta}_b[k] \tag{16}$$

is subtracted from the tilt feedback,

$$u[k] = -K_d(\hat{\theta}_b[k] - x_f[k],\, \hat{\dot{\theta}}_b[k],\, \hat{\dot{\theta}}_w[k]), \tag{17}$$

so d is absorbed by x_f at steady state, requiring no separate calibration stage. This fixed-gain design is the baseline balancing controller; Table 18 summarizes the control variables.

Table 18: Control-design variables (2D and corner-balancing 3D model).

Symbol	Description
x	State vector
A, B	Continuous-time state/input matrices (linearised model)
A_d, B_d	Discrete-time (ZOH) counterparts of A, B
K_d	LQR feedback gain
Q, R	LQR state/input weighting matrices
$J(u)$	LQR cost function
$\hat{\theta}_b, \hat{\dot{\theta}}_b$	Estimated tilt angle, rate (gyroscope/accelerometer)
$\hat{\dot{\theta}}_w$	Estimated wheel rate (encoder)
a_x, a_y	Raw accelerometer readings (tangential, radial axes)
$\theta_{b,\text{accel}}$	Instantaneous accelerometer-only tilt estimate
$\dot{\theta}_{b,\text{gyro}}$	Raw gyro rate measurement
γ	Complementary-filter weight (accelerometer term)
Δt	Control/estimation loop sample period
x_f, α	Offset-correction filter state, filter coefficient
d	Constant tilt-sensor offset
θ_{th}	Tilt threshold, jump-up $\to$ balancing switch
$K_d^{(i)}$	Scheduled gain at linearisation point i (extension)

7.2 IMU-to-Body Frame Transformation

The BMI270 cannot be mounted with its own sensor axes aligned to the body frame $\{B\}$ of Figure 7: it sits on Perfboard 1 next to the flight computer, placed for wiring and packaging rather than axis alignment (Section 6.6). Every raw gyro/accel sample is therefore rotated into $\{B\}$ by a fixed mounting matrix before it reaches the complementary filter of Equation (15) or the MEKF of Section 7.6.1.

Raw samples are first put in physical units in the sensor's own frame,

$$a_{s,i} = \frac{g_0\, \bar{a}_{\text{raw},i} - a_{\text{off},i}}{a_{\text{scale},i}}, \qquad \omega_{s,i} = \frac{\pi}{180}\, \bar{\omega}_{\text{raw},i} - \omega_{\text{bias},i}, \qquad i = 1, 2, 3, \tag{18}$$

with $g_0 = 9.806{,}65\,\text{m/s}^2$ and the bias/offset/scale constants of Table 19, fixed at calibration and not re-estimated online. The sensor-frame vectors are then rotated into the body frame by the fixed mounting matrix C_{SB},

$$\omega_b = C_{SB}\, \omega_s, \qquad a_b = C_{SB}\, a_s, \tag{19}$$

where C_{SB} is the classical Z-X-Z (3-1-3) proper Euler-angle rotation matrix built from three fixed mounting angles $(\theta_1, \theta_2, \theta_3)$, set once from the as-built mounting geometry:

$$C_{SB} = \begin{bmatrix} c_1c_3 - c_2s_1s_3 & c_3s_1 + c_1c_2s_3 & s_2s_3 \\ -c_1s_3 - c_2c_3s_1 & c_1c_2c_3 - s_1s_3 & c_3s_2 \\ s_1s_2 & -c_1s_2 & c_2 \end{bmatrix}, \qquad c_i = \cos\theta_i,\; s_i = \sin\theta_i, \tag{20}$$

with $(\theta_1, \theta_2, \theta_3) = (-30°,\; 54.7356°,\; 45°)$. The middle angle is exactly $\arccos(1/\sqrt{3})$, the angle between a cube edge and its space diagonal, consistent with a package mounted flush to a face while the axis the controller cares about at corner-up runs along the body diagonal from O. C_{SB} is verified orthonormal (unit determinant, unit row norms) whenever the mounting angles are set, guarding against a transcription error in $(\theta_1, \theta_2, \theta_3)$.

Table 19: IMU calibration constants, applied per Equation (18) before the mounting rotation of Equation (19).

Constant	Axis 1	Axis 2	Axis 3
Gyro bias ω_{bias} (rad/s)	$+0.002348$	-0.001140	-0.000762
Accel offset a_{off} (m/s^2)	-0.092282	-0.196510	$+0.050664$
Accel scale a_{scale} (–)	0.991085	0.991035	1.003787

7.3 Gain Scheduling and Mode Transitions

The controller switches from the open-loop jump-up sequence to the LQR law once $|\hat{\theta}_b| < \theta_{\text{th}}$, with hysteresis around θ_{th} to prevent mode chattering. The baseline uses a single fixed gain K_d at the upright equilibrium. If recovery degrades near the edges of the operating range, where the small-angle linearisation weakens, gain scheduling will be added: additional gains $K_d^{(i)}$ computed at several linearisation points, with the controller interpolating between them based on $\hat{\theta}_b$, reusing the same discrete structure and offset filter. The mode-transition logic gains a third mode for the 3D cube, Edge-to-Corner jump-up, triggered once edge balancing is stable, followed by a second threshold switch into the corner LQR once estimated orientation nears corner-up; the full state machine is given in Table 21.

7.4 Pure LQR Control Logic

Section 7.1 derives K_d offline, once, from the linear model. What actually executes on the flight computer every control tick is the sequence below, stated independent of how K_d was computed so that the running controller can be read, checked and re-implemented without re-deriving the gain. Table 18 (2D) and its 3D counterpart of Section 3.4 define every symbol used; the fully detailed pseudocode, including the supervisory state machine, is in Appendix D.

1. **Sense.** Read raw gyro rate $\dot{\theta}_{b,\text{gyro}}$ and accelerometer (a_x, a_y) from the IMU over SPI; read wheel angle(s) from the encoder(s) over the moteus CAN reply.
2. **Estimate tilt.** Form the accelerometer-only estimate $\theta_{b,\text{accel}}[k] = \tan^{-1}(-a_x[k]/a_y[k])$ and fuse it with gyro integration through the complementary filter of Eq. (15) to obtain $\hat{\theta}_b[k]$.
3. **Remove offset.** Update the low-pass offset state $x_f[k]$ (Eq. (16)) and subtract it from $\hat{\theta}_b[k]$.
4. **Assemble the state.** Build $\hat{x}[k] = (\hat{\theta}_b[k] - x_f[k],\ \hat{\dot{\theta}}_b[k],\ \hat{\dot{\theta}}_w[k])$ in 2D, or the 9-vector $(\theta, \omega_b, \omega_w)$ of Section 3.4 in 3D.
5. **Check the mode.** If the supervisory state machine (Table 21) is not in a balancing mode, skip to step 8 and issue that mode's own open-loop command instead of $u[k]$ below.
6. **Compute the command.** $u[k] = -K_d\,\hat{x}[k]$.
7. **Saturate and manage momentum.** Clip $u[k]$ to the driver's continuous current rating (Appendix A); if a wheel speed $\hat{\dot{\theta}}_w$ exceeds its desaturation threshold, add a small bias torque directed to bleed the stored momentum without materially perturbing $\hat{\theta}_b$.
8. **Actuate.** Issue $u[k]$ (or the mode's open-loop command) as a CAN-FD torque request to each moteus driver.
9. **Advance.** $k \leftarrow k + 1$; repeat at the next 1 kHz tick (Section 7.6.1).

This is the loop exercised in simulation (Section 7.5) before it is exercised on hardware, and it is the loop whose worst-case timing is bounded by construction in Section 7.6.1.

7.5 Simulation Environment: Simulink/Simscape Multibody

The control law of Sections 7.1–7.4 is validated in simulation before any closed-loop hardware run, on a plant built from the same CAD geometry as the physical prototypes rather than from the lumped-parameter equations alone — so that the simulation catches geometry effects (asymmetric mass distribution, off-axis clearances) the analytic model cannot represent.

CAD-to-Simscape workflow. Each CAD assembly — the 1-DoF rig of Section 5.1.2 and the 3D cube of Section 5.1 — is exported as a STEP assembly and imported into Simscape Multibody, which converts each part and mate into a rigid body and joint in a multibody tree. The pivot bearing of the 1-DoF rig and the corner/edge contact features of the 3D cube are modelled as the corresponding revolute/point-contact joints, so that the simulated degrees of freedom match Sections 3.2 and 3.4 exactly rather than approximating them. Per-part mass and inertia are taken directly from the CAD model, which is the same convention that makes the mass roll-up of Section 5.1.1 auditable: the simulation, the sizing, and the fabricated part all trace to one model.

Actuation and sensing blocks. Each motor is represented as a torque source at its joint, saturated and rate-limited to the envelope of Section 6; the brake is a discrete event that zeroes the relevant wheel-body relative velocity over the Δt of Equation (12). Simulated IMU and encoder blocks apply noise and bias models set from the sensors' characterised behaviour once measured on hardware, so that the same estimator and control logic of Section 7.4 — not a re-implementation of it — runs against simulated sensor data in a controller-in-the-loop configuration before it runs against real sensors.

Test cases. Three cases are run against the LQR baseline: recovery from an initial tilt offset, which establishes the region of attraction; rejection of an impulse disturbance applied at the body, which is the simulated counterpart of the disturbance-recovery criterion the hardware is later judged against; and behaviour under wheel-speed saturation, which exercises the momentum-dumping step of Section 7.4 and is the case most likely to be missed by a purely linear analysis, since saturation is precisely where the linear model ceases to describe the plant. Each case is re-run with the plant parameters perturbed about their nominal values, by a magnitude set once the identification tolerances of Section 3.3 are measured, so that the gains are shown not to be knife-edge: a gain set that stabilises only the nominal plant is not a result, because the nominal plant is not the one that will be built.

Direction of validation. The measured parameters of Section 3.3 are substituted into the plant and the gains re-derived before hardware balance is attempted; the simulation is validated against the measured plant, never the reverse. A discrepancy between simulated and measured closed-loop behaviour is treated as a finding about the model.

No Simulink/Simscape model or CAD export exists yet at the time of writing; this subsection states the workflow and the intended test cases, and is populated with the model, the exported geometry and the resulting response plots once built.

7.6 System Integration

7.6.1 Onboard Software Pipeline

The subsystems of Figure 11 are coordinated by a single deterministic task on the Teensy, whose structure is given in Figure 15. Each 1 ms cycle executes the same fixed sequence: sensor ingest reads the BMI270 over SPI and the moteus replies over CAN; the MEKF propagates attitude on the gyro and corrects it against the accelerometer while estimating gyro bias; the state vector $x = [g_B, \omega, h]$ is assembled from the estimator output and the wheel speeds; the LQR law of Section 7.1 evaluates $u = -Kx$ with a penalty on stored momentum h to bleed off wheel saturation; and the resulting commands leave as CAN torque requests to the drivers and a PWM signal to the brake servos. Telemetry is tapped off the assembled state to SD and Wi-Fi outside the control path, so logging cannot perturb loop timing.

The fixed ordering matters: with no dynamic allocation and no blocking calls, the worst-case cycle time is bounded by construction rather than measured after the fact, which is what makes the 1 kHz sample period of the discrete LQR design a design guarantee rather than an average.

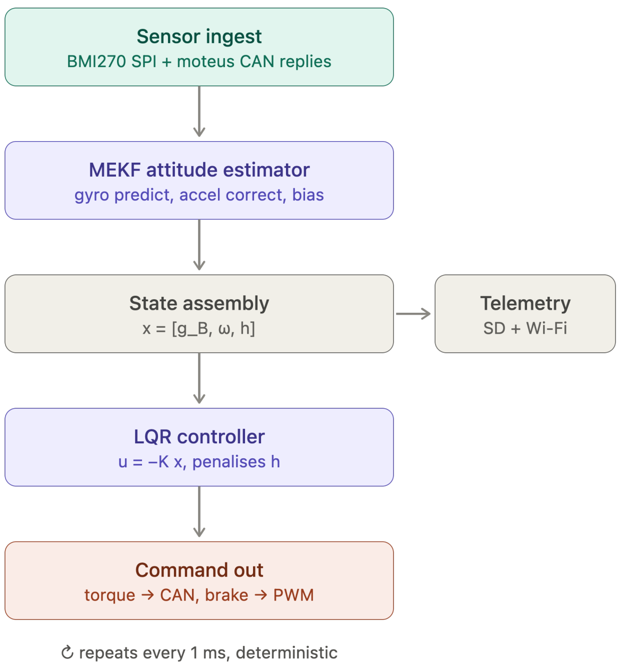

Figure 15: Teensy 4.1 control-loop software pipeline. The estimator and controller stages implement the design of Section 7.1; telemetry is a side-branch off the control path.

7.6.2 Control Software Architecture

Sections 7.1–7.4 derive the control law and its executable logic; this subsection states how that law is realised as software, which is a separate set of decisions. The implementation detail — the mode state machine, the estimator and controller in executable pseudocode, the fixed gains and tuning parameters, the timing budget and the fault responses — is collected in Appendix D, so that a re-tuned gain does not require an edit to the design sections.

The software is organised in the four layers of Table 20. The division is not stylistic: each layer has a different failure mode and a different verification method, and the interfaces between them are what allow the controller validated in simulation (Section 7.5) to be the same code that runs on the cube.

Table 20: Control software layers. Each is verified by a different method, which is why the boundaries are held.

Layer	Responsibility	Verified by
Hardware abstraction	Sensor and driver I/O, timing, bus transactions	Bench read-back; loop-timing measurement
Estimation	Attitude and rate estimate from IMU and encoders (Sec. 7.1)	Static and dynamic bench tests against a known reference
Control	State feedback, momentum management, saturation handling	Simulation on the identified model, then hardware
Supervision	Mode transitions, arming, watchdog and fault response	Fault-injection tests on the rig

Operating modes. The supervisory layer sequences the modes of Table 21, with hysteresis on every threshold to prevent chattering at a boundary, and with a disarm transition available from every mode. The modes are specified as a state machine because the critical transitions are the descending ones: a fall detected during corner balance must reach a safe state with the wheels commanded down, not merely switch to a different gain.

Table 21: Operating modes and transitions. Every mode has a disarm exit.

Mode	Behaviour	Exit condition
Idle / disarmed	Drivers de-energised; sensors and telemetry live	Explicit arm command
Wheel spin-up	Open-loop speed ramp to the jump-up momentum target	Target wheel speed reached
Jump-up	Brake actuation and momentum transfer (Sec. 4.1)	Tilt inside the balancing capture band
Edge balance	Single-axis state feedback (Sec. 7.1)	Corner-transition command, or fall detected
Corner balance	Three-axis state feedback (Sec. 3.4)	Slew command, or fall detected
Slew / rotation	Reference tracking between equilibria	Reference reached; revert to balance
Fault / safe	Wheels commanded down, drivers disarmed	Manual reset

Real-time guarantees. The loop-rate claim of Section 7.1 is a design guarantee only if the worst-case cycle is bounded by construction, which is why the pipeline of Figure 15 uses a fixed stage ordering with no dynamic allocation and no blocking calls, and why telemetry is a side-branch rather than a stage.

Fault handling. The failure cases the software must survive are sensor dropout, bus error, wheel saturation and loss of the balancing equilibrium. The first three have defined responses that keep the machine controlled; the fourth is a controlled shutdown rather than a recovery attempt, since a cube that has left its capture region cannot be recovered by the linear law and continued actuation only adds energy to the fall.

8 Results

8.1 Validation Results to Date

This subsection reports evidence against the acceptance criteria of Table 1, in the order the stages are entered. Where no result exists yet, the criterion is stated so the gap is explicit rather than silent.

8.1.1 1-DoF Control Validation

The single-axis rig is the protected experimental item of Section 2.3, and its balance result is the evidence carried into this document. Gate G3 requires the rig to balance for at least 60 s unassisted.

Beyond the duration criterion, three quantities are recorded because they determine whether the result transfers to the integrated cube rather than merely holding on the rig. Recovery from a defined impulse establishes that the balance is a stabilised equilibrium and not a fortunate initial condition. Steady-state wheel-speed drift measures momentum accumulation: a controller that balances while monotonically spinning up is one that will saturate, and the drift rate predicts when. Control effort is compared against the continuous and peak current limits of the driver given in Appendix A, since a balance sustained only by exceeding the continuous rating is not a balance that survives to the demonstration.

Balancing performance examples. Table 22 reports representative runs against Gate G3, each initialised from rest at the upright equilibrium with no applied disturbance.

Table 22: Representative unassisted balance runs on the single-axis rig, against the Gate G3 threshold of 60 s.

Run	Balance duration	Steady-state wheel-speed drift	Peak driver current
1	> 60 s	–	0.36 A
2	> 60 s	–	0.38 A
3	> 60 s	–	0.40 A

Note: No constant drift was appreciated during static zero-position balancing. When the center of mass was uncalibrated, small oscillations were observed instead of steady drift due to torque saturation.

Recovery-angle disturbance tests. Beyond the zero-disturbance case, the rig was displaced by hand to a defined initial angle, measured from the support, and released to test whether the controller recovers to upright rather than merely holding a small perturbation. Table 23 reports the three angles used to characterise this margin.

Table 23: Recovery outcome versus the initial displacement angle, measured from the support.

Initial angle	Recovered?	Recovery time	Peak driver current
7°	Yes	1.2 s	0.86 A
8°	Yes	1.5 s	0.98 A
12°	Yes	2.1 s	1.34 A

The largest initial angle tested to date, again measured from the support, is 24°; recovery from this angle was also satisfactory, so the demonstrated disturbance-rejection envelope currently extends at least this far beyond the three angles of Table 23. This angle represents the physical hard stop of the experimental rig support as well as the maximum tilt evaluated in simulation, fully satisfying the requirements defined in the testing plan.

8.1.2 3-DoF Control Validation

For the 3D model, control-loop validation to date has achieved edge balance on the final prototype using the 1-DoF control law; corner balance has not yet been demonstrated. Table 24 sets out the result template that will be populated once edge-balance testing on the cube is complete, mirroring the 1-DoF quantities of Section 8.1.1 so the two are directly comparable, and Table 25 mirrors the recovery-angle protocol of Table 23 for the cube. TODO: populate both tables from edge-balance testing on the final prototype, run the corner-balance case once edge balance is closed out, and compare against the acceptance criteria of Table 1 and the 1-DoF bench results of Section 8.1.1.

Table 24: Result template: unassisted balance on the cube, edge and corner equilibria, against the Gate G3 threshold of 60 s.

Equilibrium	Balance duration	Steady-state wheel-speed drift	Peak driver current
Edge	TBD	TBD	TBD
Corner	TBD	TBD	TBD

Table 25: Result template: recovery from a defined initial angle on the cube, measured from the support, edge equilibrium.

Initial angle	Recovered?	Recovery time	Peak driver current
7°	TBD	TBD	TBD
8°	TBD	TBD	TBD
12°	TBD	TBD	TBD
24°	TBD	TBD	TBD

9 Conclusions & Next Steps

9.1 Conclusions

The team has demonstrated the ability to build and control the Cubli across the planned difficulty progression: a working 1-DoF model, and edge balance on the final 3D prototype, both achieved through the staged validation approach described in this document. Several of the originally planned tasks have been rescheduled in response to deadline pressure and issues encountered along the way, but the team has continued to make progress toward the final objective. Corner balance remains the outstanding milestone and is expected to be resolved in the days before final delivery and presentation. Cubli has confirmed itself to be a tightly coupled, multidisciplinary system: reliably building and controlling it depends on close coordination and continuous communication across the team's subsystem interfaces.

9.2 AI-Driven Control & Dynamic Coefficient Tuning

9.2.1 Motivation: Limitations of a Fixed Gain

The balancing controller of Section 7.1 computes a single gain K_d offline, from an (A, B) pair built from the parameters of Tables 6 and 8. That gain is optimal for exactly one plant: the one those numbers describe. It is retained here as the baseline precisely because it is simple and verifiable, but its optimality is conditional in three ways that the hardware of this project is expected to violate.

- **Parametric drift.** C_b and C_w are not obtained from CAD and are carried as estimates until the identification tests of Table 7. Bearing friction, wiring drag and motor cogging also change with temperature, wear and reassembly, so the identified values are a snapshot rather than a constant. The inertia tensor I and the centre-of-mass offset r_{cm} are subject to the same objection at the few-percent level, and r_{cm} in particular shifts whenever the pack or harness is disturbed.
- **Payload and configuration changes.** Any added mass changes m, r_{cm} and I simultaneously, and therefore changes A. A gain tuned for one configuration is not merely suboptimal for another; near a fully unstable equilibrium it may be destabilising.
- **Contact and surface variation.** The pivot is modelled as an ideal frictionless point or line. A real corner resting on a real surface slips, deforms and dissipates energy in a way that is neither in the model nor constant across surfaces, and this is the dominant unmodelled effect for the walking objective of Section 1.3.

The standard remedy within the existing design is gain scheduling on $\hat{\theta}_b$ (Section 7.1), but scheduling interpolates between gains that were themselves computed from the same assumed parameters. It addresses the weakening of the small-angle linearisation; it does not address a plant that differs from the model. The proposal here is to keep the verified feedback law and adapt its coefficients online.

9.2.2 Proposed Hybrid Architecture

The intended structure is deliberately conservative: a classical inner loop whose stability properties are understood, wrapped by a slow outer loop that adjusts the inner loop's coefficients. Writing ϕ for the tunable coefficient vector,

$$u[k] = -K_d(\phi[k])\,\hat{x}[k], \qquad \phi[k+1] = \phi[k] + \Delta\phi\big(\hat{x}[k],\, \hat{x}[k-1], \dots\big), \tag{21}$$

where $\hat{x}$ is the estimated state of Section 7.1 and $\Delta\phi$ is produced by the tuner. Two properties of (21) matter more than the choice of learning algorithm. First, the inner loop remains a linear state feedback at every instant, so the controller that runs on the hardware is the one already implemented and tested. Second, ϕ is updated on a timescale far slower than the control period Δt, which keeps the frozen-gain analysis approximately valid and prevents the tuner from acting as an unmodelled fast feedback path. Three candidate tuners are of interest, and they are not equivalent in cost:

- **Bayesian optimisation** over a low-dimensional parameterization of Q and R, minimising a measured cost from short balancing trials [7]. This is the cheapest option and the one with a direct precedent on this class of hardware: automatic LQR tuning has been demonstrated on an inverted pendulum with on the order of tens of experiments [8], and a self-tuning LQR has been run on the Cubli's own ancestor platform at ETH [9]. It tunes between trials, not within them.
- **Adaptive dynamic programming**, which solves the LQR problem online from measured data rather than from (A, B), converging toward the optimal gain for the true plant without an explicit model [10]. This is the option with the strongest theoretical guarantees and the strictest excitation requirements.
- **Reinforcement learning**, in which a policy trained in simulation outputs $\Delta\phi$ continuously from the state-error history. This is the most capable and the least verifiable, and is treated below as the exploratory branch rather than the default.

Deployment reservations. The literature supporting learned control is overwhelmingly validated in simulation or on well-instrumented hardware, and the failure mode of an online tuner on an open-loop-unstable plant is a fall, not a degraded score. Any deployment on this cube should therefore be gated behind: a projection of ϕ onto a set for which the frozen-gain closed loop is verified stable across the parameter uncertainty; a fallback to the fixed K_d of Section 7.1 on any estimator fault or divergence; and bench trials with the cube tethered. This is future work in the literal sense — none of it is a committed deliverable of the present design, and the baseline controller is not contingent on any of it.

9.2.3 Module 1: Dynamic Gain Tuning

The tuner does not emit motor currents. It emits the coefficients from which K_d is computed, which bounds what a bad update can do. Two parameterizations are worth comparing:

- **Weight-space tuning.** The policy outputs the diagonal entries of the cost matrices, $\phi = (\log q_1, \ldots, \log q_9,\ \log r_1, \log r_2, \log r_3)$ with $Q = \mathrm{diag}(q_i) \succ 0$ and $R = \mathrm{diag}(r_i) \succ 0$, and K_d follows from the discrete algebraic Riccati equation. The logarithmic parameterization enforces positive definiteness by construction, so every ϕ the policy can emit corresponds to a well-posed LQR problem and, for a stabilisable (A_d, B_d), to a stabilising gain for the *modelled* plant. The cost is that a Riccati solve must run onboard whenever ϕ changes, which is what confines updates to a slow outer loop.
- **Gain-space tuning.** The policy perturbs the entries of K_d directly, or equivalently the K_p, K_d pair of a PD form, avoiding the Riccati solve at the price of losing the structural guarantee above. This is the practical choice if the outer loop must run fast, and it makes the projection step of the previous paragraph mandatory rather than merely prudent.

For the 3D corner-balancing case the state is $x = (\theta, \omega_b, \omega_w) \in \mathbb{R}^9$ and K_d is 3×9 (Section 3.4), so tuning the 54 free entries directly is not sensible on the available data. Weight-space tuning reduces the problem to 12 coefficients, and symmetry of the three wheel axes reduces it further if the wheels are treated as identical. A reward or cost of the form

$$c[k] = \hat{x}^{\mathrm{T}}[k]\, W_x\, \hat{x}[k] + \rho\, \|u[k]\|^2 + \sigma\, \|\omega_w[k]\|^2 \tag{22}$$

Equation (22) adds an explicit penalty on stored wheel momentum, which the nominal LQR cost $J(u)$ of Section 7.1 does not distinguish from any other state deviation. This term is the one directly relevant to the saturation problem raised in Section 1.2: it expresses a preference for holding attitude with the wheels near zero speed, leaving momentum headroom for disturbance rejection.

9.2.4 Module 2: Sim-to-Real Transfer via Domain Randomization

A policy trained against a single nominal parameter set learns to exploit that parameter set. Domain randomization instead samples the plant parameters from a distribution at the start of each training episode [11], so that the policy is optimised against a family of plants rather than one, and the transfer method has a substantial record on physical legged hardware [12]. For this cube the randomized quantities follow directly from the identification gaps already documented:

- **Damping:** C_b, C_w over a wide interval, since these are estimated rather than measured until the identification of Section 3.3 closes.
- **Inertia and mass distribution:** I, m and r_{cm} perturbed by a few percent, covering print-density variation, fastener placement and harness routing.
- **Actuation:** K_m within motor tolerance, plus a current saturation limit and a commanded-torque delay, so the policy cannot assume an ideal actuator.
- **Contact:** pivot friction and restitution, and a small random ground slope — the term standing in for the surface variation that motivates this module.
- **Sensing:** gyro bias walk, accelerometer noise and an injected vibration component, matching the estimator degradation identified in Section 1.1 as the practical limit on balancing.

The sensing term requires particular care. The accelerometer corruption on this platform is generated by the controller's own wheel activity, so it is correlated with the policy's actions rather than independent of them. A randomization scheme that injects only white sensor noise will understate the difficulty, and a policy that transfers under it should not be assumed to transfer on hardware. Whether the randomized distribution covers the real cube is an empirical question, settled by bench testing; the principal failure mode of this module is a distribution selected for convenience rather than representativeness.

9.2.5 Module 3: Jump-Up Energy Optimisation

The jump-up analysis of Section 4.1.1 treats the manoeuvre as a spin-up to $\omega_{\max}$ followed by an abrupt stop, with all deviation from the ideal absorbed into a single brake-amplification factor $\beta = \tau_b/(\tau_b - \tau_g)$. That model is adequate for sizing and is deliberately conservative, but it is a poor description of what the mechanism does: the servo-driven barrier has finite travel, the collision is not instantaneous, and the resulting torque profile is neither a step nor known in advance.

The proposal is to treat the torque command over the manoeuvre window as a profile to be optimised rather than a constant. Parameterizing the pre-brake and post-impact torque commands by a small vector ψ — for instance the timing of brake release relative to the wheel-speed trajectory, and the shape of the recovery torque applied by the remaining wheels immediately after impact — the objective is

$$\min_{\psi} \;\; \mathbb{E}\Big[\, \|\theta(t_f) - \theta^\star\|^2 + \mu\, \|\omega_b(t_f)\|^2 \Big] \quad \text{subject to} \quad \omega_w \le \omega_{\max}, \;\; |u_i| \le u_{\max}, \tag{23}$$

where t_f is the instant the controller hands over to the balancing law at $|\hat{\theta}_b| < \theta_{\text{th}}$. Minimising residual body rate at handover is the quantity of interest, because a jump that reaches the equilibrium with excess rate demands recovery authority the balancing controller may not have. Because each evaluation of (23) is a physical jump, the sample budget is small and the appropriate method is Bayesian optimisation rather than a policy-gradient method, with the initial design taken from the analytical h_w of Equation (10).

Two constraints on this module must be stated. First, it presupposes the measurement described in Section 4.1.1: until the spin-down test reports both the impulse-equivalent mean $\bar{\tau}_b = h_w/\Delta t$ and the peak torque, there is no validated model to optimise against and no way to distinguish an improved profile from a mis-measured one. Second, the structural caveat carries over unchanged — optimisation must not be permitted to search toward shorter stopping times, since τ_b is simultaneously a load on the barrier, bolt, bracket and spokes, and the sizing value was chosen for the dynamics and not for that load path. The search space must be bounded by the structural case, not by the dynamic one.

9.2.6 Relevance Beyond the Bench

Two application claims follow, and both are narrower than the general enthusiasm for learned control would suggest.

For **CubeSat attitude control**, the relevant fact is that a spacecraft's inertia tensor is known imprecisely at integration and changes in flight — propellant depletion, deployable appendages, thermal distortion — while reaction-wheel friction rises with bearing age. These are the orbital analogues of the parametric drift of Section 7.1, and they are conventionally handled by conservative fixed gains with margin, at a cost in pointing performance. A tuner that adapts Q and R against measured performance addresses exactly this gap. The qualification barrier for flying learned control is high and unaddressed here; the architecture of (21) is chosen partly because a bounded outer loop over a classical inner loop is the form most likely to survive that scrutiny.

For **low-gravity surface mobility**, the case is stronger, because the plant is genuinely unknown before arrival. Regolith mechanical properties on a small body are not characterised to the precision a hopping controller would want, and the flight-proven hoppers cited in Section 1.2 [4, 5] move ballistically rather than under closed-loop control for this reason. The contact randomization of Module 2 and the profile optimisation of Module 3 map onto that problem directly: a hop whose landing attitude is controlled, on a surface whose properties were sampled from a distribution rather than assumed, is the capability that separates directional walking from a ballistic tumble [6]. The cube tests this under 1 g, where the destabilising torque is larger and never relents — a harder case in that one respect, as argued in Section 1.2, though it does not reproduce the low-gravity contact dynamics that would decide the real manoeuvre.

References

[1] Mohanarajah Gajamohan, Michael Merz, Igor Thommen, and Raffaello D'Andrea. "The Cubli: A Cube that can Jump Up and Balance". In: *2012 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Vilamoura-Algarve, Portugal, Oct. 2012, pp. 3722–3727. DOI: 10.1109/IROS.2012.6385896.

[2] Mohanarajah Gajamohan, Michael Muehlebach, Thomas Widmer, and Raffaello D'Andrea. "The Cubli: A Reaction Wheel Based 3D Inverted Pendulum". In: *2013 European Control Conference (ECC)*. Zurich, Switzerland, July 2013, pp. 268–274.

[3] Michael Muehlebach and Raffaello D'Andrea. "Nonlinear Analysis and Control of a Reaction-Wheel-Based 3-D Inverted Pendulum". In: *IEEE Transactions on Control Systems Technology* 25.1 (Jan. 2017), pp. 235–246. DOI: 10.1109/TCST.2016.2549544.

[4] Tetsuo Yoshimitsu, Takashi Kubota, and Ichiro Nakatani. "New Mobility System for Small Planetary Body Exploration". In: *Proceedings of the 1999 IEEE International Conference on Robotics and Automation (ICRA)*. 1999, pp. 1404–1409.

[5] Tra-Mi Ho, Vladimir Baturkin, Christian Grimm, Jan Thimo Grundmann, et al. "MASCOT—The Mobile Asteroid Surface Scout Onboard the Hayabusa2 Mission". In: *Space Science Reviews* 208 (2017), pp. 339–374. DOI: 10.1007/s11214-016-0251-6.

[6] Alexander Spiridonov et al. "SpaceHopper: A Small-Scale Legged Robot for Exploring Low-Gravity Celestial Bodies". In: *2024 IEEE International Conference on Robotics and Automation (ICRA)*. arXiv:2403.02831. 2024, pp. 3464–3470.

[7] Bobak Shahriari, Kevin Swersky, Ziyu Wang, Ryan P. Adams, and Nando de Freitas. "Taking the Human Out of the Loop: A Review of Bayesian Optimization". In: *Proceedings of the IEEE* 104.1 (2016), pp. 148–175. DOI: 10.1109/JPROC.2015.2494218.

[8] Alonso Marco, Philipp Hennig, Jeannette Bohg, Stefan Schaal, and Sebastian Trimpe. "Automatic LQR Tuning Based on Bayesian Optimization with Ninety-Nine Percent Confidence". In: (2016), pp. 1557–1563. DOI: 10.1109/ICRA.2016.7487292.

[9] Sebastian Trimpe, Alexander Millane, Simon Doessegger, and Raffaello D'Andrea. "A Self-Tuning LQR Approach Demonstrated on an Inverted Pendulum". In: *IFAC Proceedings Volumes (19th IFAC World Congress)* 47.3 (2014), pp. 11281–11287. DOI: 10.3182/20140824-6-ZA-1003.01455.

[10] Frank L. Lewis, Draguna Vrabie, and Kyriakos G. Vamvoudakis. "Reinforcement Learning and Feedback Control: Using Natural Decision Methods to Design Optimal Adaptive Controllers". In: *IEEE Control Systems Magazine* 32.6 (2012), pp. 76–105. DOI: 10.1109/MCS.2012.2214134.

[11] Josh Tobin, Rachel Fong, Alex Ray, Jonas Schneider, Wojciech Zaremba, and Pieter Abbeel. "Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World". In: *2017 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. 2017, pp. 23–30. DOI: 10.1109/IROS.2017.8202133.

[12] Jemin Hwangbo et al. "Learning Agile and Dynamic Motor Skills for Legged Robots". In: *Science Robotics* 4.26 (2019), eaau5872. DOI: 10.1126/scirobotics.aau5872.

[13] Matthias Hofer, Michael Muehlebach, and Raffaello D'Andrea. "The One-Wheel Cubli: A 3D Inverted Pendulum That Can Balance with a Single Reaction Wheel". In: *Mechatronics* 91 (2023), p. 102965. DOI: 10.1016/j.mechatronics.2023.102965.

A Bill of Materials

The components supplied for this project constitute an *electronics* bill of materials: the actuation chain, wheel and attitude sensing, on-board compute, communications and power distribution are specified, but no structural or mechanical hardware is. The reaction wheels, the frame, the motor mounts, the fasteners and every printed bracket are the team's design responsibility, and are treated as such in Section 5. Table 26 lists the supplied items grouped by subsystem. The actuation figures derived from these parts — the torque constant, the achievable torque and wheel speed, and the inertia that sizes the wheel — are developed where they are used, in Sections 4.1.1 and 4.3.

The architecture the BOM implies is a two-level control stack: each mjbots moteus-n1 closes a high-rate field-oriented current loop on its motor, using the MA600 encoder for rotor angle, while the Teensy 4.1 runs the outer attitude-estimation and state-feedback loop over CAN-FD — the same inner-tachometer / outer-attitude layering used in spacecraft reaction-wheel assemblies. A mechanical wheel brake, following the original ETH Cubli [2], provides the near-impulsive momentum transfer required for the jump-up.

Procurement is the single largest schedule risk. The long-lead items are not electronic but mechanical: the frame stock and the printed structure gate the entire build, and with three actuation sets and no spare, the loss of a motor or driver in service would force the reduced-actuation one-wheel variant [13]. These items are ordered first. Consumables and bench equipment not covered by the supplied set — wire, connectors, fasteners, filament, the current-limited bench supply and the battery-safety kit — are tracked separately in the build documentation.

Table 26: Supplied bill of materials, grouped by subsystem.

Ref	Item	Qty[a]	Key specifications	Function
Actuation				
1	T-Motor MN4006 KV380	3	18N24P outrunner; 380 rpm/V; 68 g; ⌀44.35 mm × 21 mm; 194 mΩ; peak 16 A, 380 W	Reaction-wheel torque source
2	mjbots moteus-n1	3	10–54 V; 100 A peak / 9 A cont.; 15–30 kHz FOC; 5 Mbit/s CAN-FD; 14.6 g	FOC current/torque driver (inner loop)
11	TowerPro MG92B servo	3	metal-gear micro servo; ≈0.29 N m (≈3 kg cm) at 6 V; 50 Hz PWM	Wheel brake actuator (jump-up)
Sensing				
3	mjbots MA600 breakout + ring magnet	3	TMR absolute angle; SPI 6 MHz, 3.3 V; 16-bit (65,536 counts/rev); air-gap ≤ 0.5 mm	Wheel angle: commutation + state feedback
10	SparkFun BMI270 6-DoF IMU	1	3-axis accel + 3-axis gyro (no magnetometer); I^2C 400 kHz (Qwiic); ODR up to ≈ 1 kHz	Body attitude reference (tilt + rates)
Compute and Communications				
4	Teensy 4.1	1	600 MHz Cortex-M7 + FPU; native CAN-FD (CAN3); 5 V supply	Flight computer: estimator + control law
5	CAN-FD adapter (Teensy 4.1)	1	CAN-FD transceiver, logic-level → differential; 5 Mbit/s	Bus transceiver for the MCU

Continued on next page

Table 26 (continued)

Ref	Item	Qty[a]	Key specifications	Function
6	mjbots mjcanfd-usb-1x	1	USB ↔ CAN-FD, 5 Mbit/s	Bench interface (calibration, tuning)
7	JST-PH3 CAN cables	≥4	2 A AC/DC; 3-conductor daisy-chain	CAN bus interconnect
8	Seeed XIAO ESP32-C6 + antenna	1	Wi-Fi; 3.3 V UART to MCU; TX burst ≈ 300 mA	Wi-Fi telemetry / command link
9	60.4 Ω 0.5 W resistor	4	Split termination, 2 × 60.4 Ω ≈ 120 Ω bus impedance	CAN bus termination
17	4.7 nF capacitor	2	common-mode filter, bus midpoint to ground	CAN split-termination filter
Power				
12	Tattu R-Line V5.0 6S LiPo (XT60)	1	22.2 V; 1550 mA h; 150C; ≈ 35 W h; ≈ 254 g	Primary energy store
13	LM2596 buck module	1	step-down to 5.0 V; rated 2–3 A (derates without heatsink)	5 V logic rail
14	XT60 connector pair	1–2	rated ≈ 60 A	Battery power connector
15	100 µF 16 V electrolytic	2–4	16 V part — 5 V rail only	5 V-rail bulk reservoir
19	1N5819 Schottky diode	1–2	1 A, 40 V	5 V-rail back-feed protection
Integration				
16	100 nF 50 V capacitor	6–10	local decoupling at each IC	Per-IC supply decoupling
18	Double-sided perfboard	1	15 cm × 9 cm, double-sided	Power / signal distribution board

[a] Quantity per assembled cube; motors, drivers, encoders and brake servos are 3 each — one per wheel axis.

B CAD and Construction Detail

This appendix holds the mechanical detail deliberately kept out of Section 5.1: the model tree, the drawing set, the fabrication settings and the assembly sequence. The rationale for the architecture — why the frame is split, what fixes the envelope, how the wheel is tuned — stays in Section 5.1 and is not repeated here. Anything in this appendix may change without the design intent changing; anything in Section 5.1 may not.

B.1 Design Parameter Table

Table 27 is the single source of the driving dimensions referenced by every part in the model, per convention 1 of Section 5.1.1. It is the CAD-side mirror of the envelope closure of Section 4.2 and is updated on each iteration.

Table 27: Driving CAD parameters. Every dependent feature references these; no part carries a duplicated hard-coded value.

Symbol	Parameter	Value
L	Cube edge length (outer)	15 cm
t_w	Frame wall thickness	5 mm
D_w	Wheel outer diameter (Sec. 4.3)	120 mm
t_r	Wheel axial thickness (Sec. 4.3)	5 mm
d_m	Motor axis offset from cube centre	TBD
h_m	Motor mounting face height on its axis	TBD
c_w	Wheel-to-subframe axial clearance	TBD
c_o	Clearance between orthogonal wheel modules	TBD
$V_{\min}$	Minimum internal clear volume	TBD
r_c	Corner contact feature radius	TBD

B.2 Model Tree and Part List

Table 28 lists the modelled parts by assembly, with their fabrication route and current status. It is the mechanical counterpart to the electronics bill of materials of Appendix A: the parts here are the team's design responsibility and are not supplied.

B.3 1-DoF Prototype Construction

The general-assembly drawing of the 1-DoF rig is given in Figure 16. This is the final released design of the single-axis prototype: the sheet carries the orthographic set — front and side elevations of the rig on its base, and a plan and edge view of the reaction-wheel assembly — from which the plate geometry, the pivot bearing seat, the motor bolt circle and the hard-stop placement are taken. The rig as built from this drawing is shown in Figure 9.

Table 28: CAD model tree and part list. Quantities are per assembled cube unless noted; the 1-DoF rig parts are listed separately.

Ref	Part	Qty	Fabrication route	Status
Inner structure				
M1	Motor subframe	TBD	TBD	TBD
M2	Motor mount (×axis)	3	TBD	TBD
M3	Battery cradle / retention	1	TBD	TBD
M4	Electronics tray	1	TBD	TBD
M5	Brake servo bracket	3	TBD	TBD
Reaction wheel assembly				
W1	Ballasted ring (Sec. 4.3)	3	TBD	TBD
W2	Wheel hub / rotor adapter	3	TBD	TBD
W3	Ballast bolt and nut set (M6, 3×6)	18	Standard hardware	TBD
W4	Brake strike feature	3	TBD	TBD
Outer frame				
F1	Frame member / edge rail	TBD	TBD	TBD
F2	Corner contact feature	8	TBD	TBD
F3	Face panel / wheel guard	TBD	TBD	TBD
1-DoF rig (Sec. 5.1.2)				
R1	Face plate	1	TBD	TBD
R2	Pivot bearing mount	1	TBD	TBD
R3	Rig motor mount	1	TBD	TBD
R4	Hard stops / containment	TBD	TBD	TBD

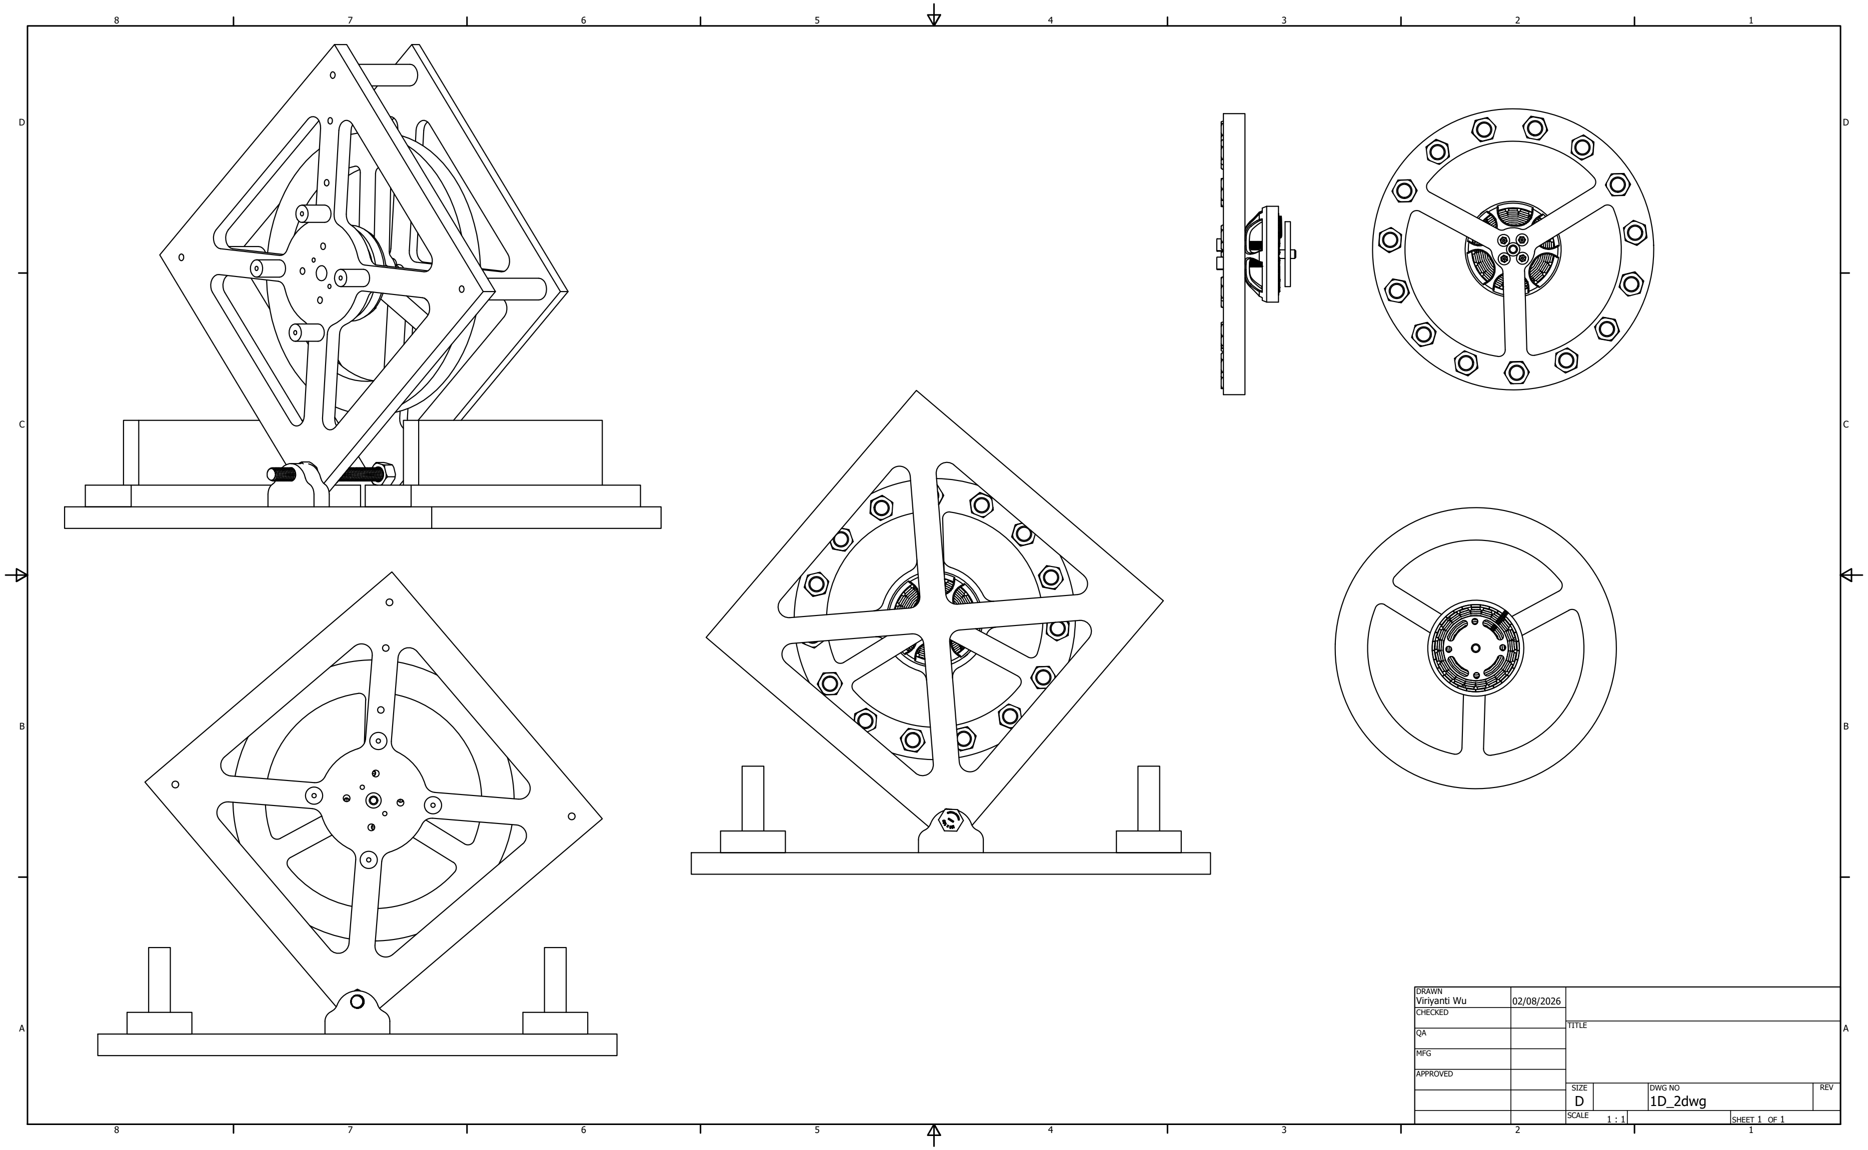

Figure 16: General-assembly drawing of the 1-DoF prototype rig — the final released design of the single-axis article — sheet 1 of 1 at 1:1. Upper left and lower left: front elevation of the rig mounted on its base, showing the diagonal plate orientation, the corner pivot and the two hard stops. Centre: side elevation through the wheel axis. Right: plan and edge views of the reaction-wheel and motor assembly, giving the spoke pattern and the ballast bolt circle used for the balancing adjustment of Section 4.3.

B.4 3D Cube Construction

The general-assembly drawing of the three-axis cube is given in Figure 17. The frame realises the two-part architecture of Section 5.1: an outer frame of edge rails and corner contact features encloses an inner structure carrying the three orthogonal reaction-wheel modules and the electronics tray. The sheet carries an isometric cutaway of the inner structure with the individual boards called out (upper left), an isometric view of the closed assembly with the three wheel modules visible through the open frame faces and the ballast bolt set annotated (upper centre), a front elevation of a single motor/wheel module in section, with the IMU location called out (right), and two front orthographic views of individual face modules giving the motor bolt circle and corner fastener pattern (lower left and centre). This drawing supersedes the concept prototype of Figure 10, which validated the two-part architecture and internal packaging but did not carry final dimensions; the edge length L and wheel geometry realised here follow the closure of Section 4.3.

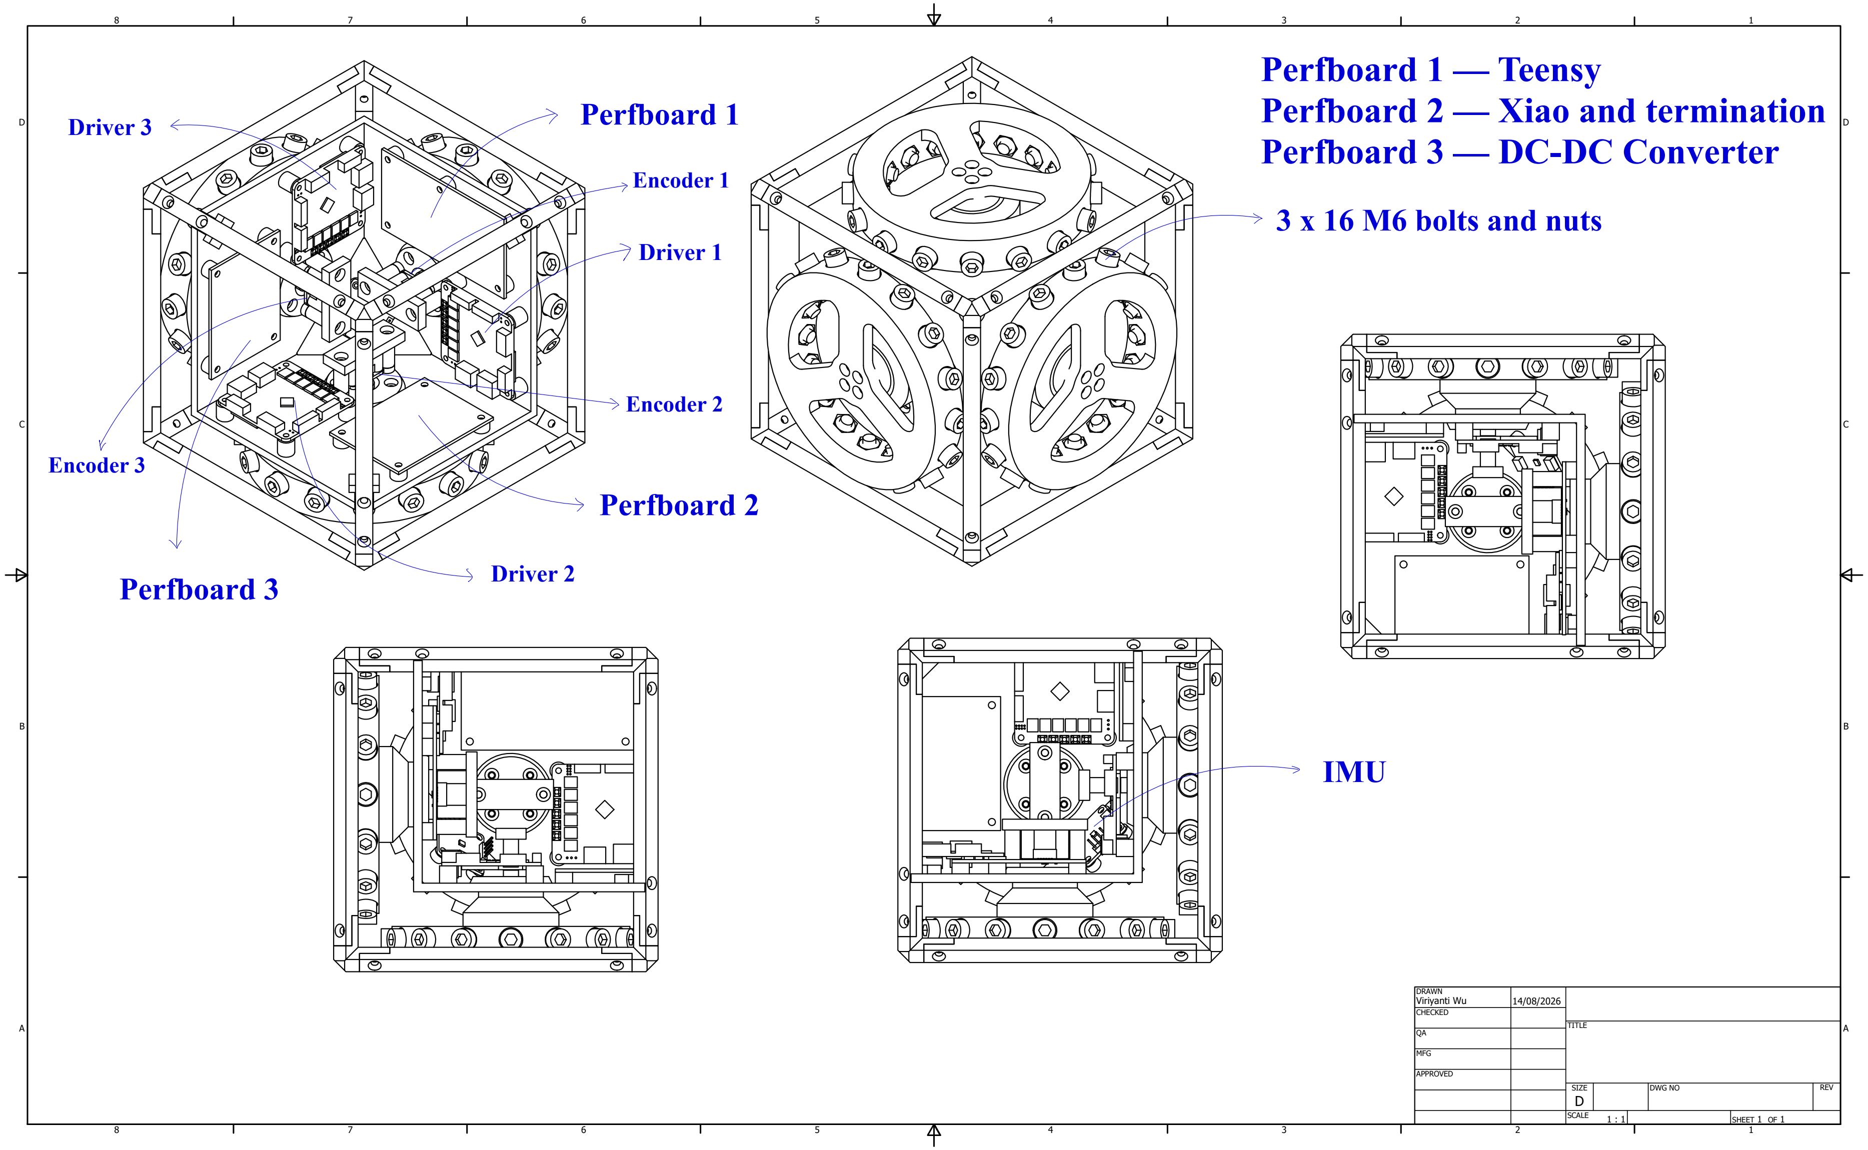

Figure 17: General-assembly drawing of the three-axis cube — the final released design of the flight article — sheet 1 of 1 at 1:1, with the electronics package called out per Section 6.6. Upper left: isometric cutaway of the inner structure, labelling the three motor drivers (Driver 1–3), the three encoder boards (Encoder 1–3) and the three perfboard assemblies (Perfboard 1: Teensy flight controller; Perfboard 2: Xiao and termination; Perfboard 3: DC-DC converter). Upper centre: isometric view of the closed assembly, showing the three orthogonal reaction-wheel modules through the open frame faces and the 3×6 M6 ballast bolt-and-nut set. Right: front elevation in section of a single motor/wheel module, locating the IMU on the inner structure. Lower left and lower centre: front orthographic views of two face modules, giving the motor bolt circle and the corner fastener pattern of Table 15.

The IMU callout above marks its physical location only: its sensing axes are not aligned with the body frame $\{B\}$ of Figure 7, and the fixed correction applied in software is given in Section 7.2.

B.5 Fabrication and Assembly

B.5.1 Print and Manufacturing Settings

The printed parts are structural, not cosmetic: the wheel ring carries centrifugal load at the ballast holes and the motor mounts carry the brake impulse, which is the design load case identified in Section 1.1. Settings are therefore recorded per part in Table 29 and treated as part of the design, since a part reprinted at a different orientation or infill is not the part that was analysed.

Table 29: Fabrication settings for the structural printed parts. Layer orientation is specified relative to the principal load direction.

Ref	Material	Infill / walls	Orientation and rationale
W1	TBD	TBD	TBD
W2	TBD	TBD	TBD
M1	TBD	TBD	TBD
M2	TBD	TBD	TBD
F1	TBD	TBD	TBD
F2	TBD	TBD	TBD

B.5.2 Assembly Sequence

The sequence of Table 30 is ordered by access: each step is one whose verification becomes impossible or destructive once a later step is complete. The encoder air gap in particular is re-checked after mechanical assembly, since it can shift when the structure is bolted up (Table 17).

Table 30: Assembly sequence. Ordering is set by access: each check must be completed before the step that encloses it.

Step	Operation	Check before proceeding
1	Wheel assembly: ring, hub, ballast bolts to the selected count	Balance check; bolt retention per Sec. 4.3
2	Motor to mount, encoder magnet to rotor	Encoder air gap and centring to the assembly spec
3	Motor modules to inner structure, three axes	Axis orthogonality; wheel clearance c_w, c_o
4	Battery and electronics tray into inner structure	Pack centring; board keep-out clearances
5	Harness install and dressing	Continuity and isolation (Sec. 6.7)
6	Inner structure into outer frame	Encoder air gaps re-checked after bolt-up
7	Contact features and guards	Contact geometry; wheel guard clearance
8	Mass properties measurement	Measured mass against Table 13

B.6 Mass Properties Verification

The CAD roll-up is an estimate until the assembly is weighed. Step 8 above closes the loop of Section 4.2: the measured mass is compared against the modelled value, and any discrepancy is carried back into $I_{w,\text{target}}$ and absorbed, where possible, by the ballast trim of Section 4.3 rather than by a re-design.

Table 31: Modelled versus measured mass properties. A discrepancy here is absorbed by ballast trim where the margin allows.

Quantity	CAD value	Measured	Discrepancy
Total cube mass M	TBD	TBD	TBD
Wheel mass m_w (each)	TBD	TBD	TBD
Wheel inertia I_w (each)	TBD	TBD	TBD
CoM offset from cube centre	TBD	TBD	TBD

C Electrical System Detail

This appendix holds the electrical detail kept out of Section 6: the full schematic capture, the connector and pin assignments, the harness drawings, the board layout and the bring-up checklists. The architectural decisions — two rails behind one E-stop, linear bus topology, perfboard rather than fabricated board — remain in Section 6. The part specifications are in Appendix A and are not repeated.

C.1 Schematic Capture

The system is captured on one flat KiCad sheet, reproduced at page size in Figure 18 and listed by functional block in Table 17. That capture is the working drawing the build is wired from. The blocks it carries are: power entry and protection (pack connector, anti-spark inrush path, combined E-stop and fuse, split to the motor bus and the logic rail); the 5 V logic rail off an LM2596 converter with back-feed protection and per-device decoupling; the CAN-FD bus as a linear chain across the three moteus drivers and the flight computer, terminated at the two physical ends only; the per-axis encoder SPI interfaces; the Teensy 4.1 flight computer with the BMI270 IMU; the ESP32 telemetry link; and the three brake servo drives.

Note: for the provided N35 1/4" $\times$ 1/8" diametric disk magnet, the optimal distance in order to achieve the 20 mT – 80 mT nominal magnetic field for the MA600 encoder is 3.3 mm.

Each block of the capture is checked against its own layout, mounting and connection constraint before bring-up. That constraint table (Table 17) is given in Section 6.2, ahead of the schematic itself, and is not repeated here. Component current ratings are given in Appendix C.2 (Table 32).

C.2 Component Current Ratings

Table 32 states the current each component on the 5 V logic rail and the pack rail is rated to draw. It replaces a measured-consumption table: at this stage no bench data exists, and a table of datasheet ratings is a stronger design input than an empty one, since the 5 V rail is sized against the coincident worst case — telemetry transmit burst, servo stall and full sensor load at once — and not the sum of typical values. Two entries are marked as unavailable rather than estimated: neither manufacturer publishes a stall or supply-current figure for that part, and a fabricated number would be worse than an honest gap. Both are closed by the bench measurement of Table 34 during electrical bring-up (Section 6.7).

Table 32: Component current ratings, by rail. Datasheet figures where published; measured values replace these per Table 34.

Component	Rail	Rated current	Basis
Teensy 4.1	5 V	≈ 100 mA	Board only, 600 MHz, no shields (PJRC community data)
XIAO ESP32-C6 telemetry	5 V	≈ 300 mA	TX burst (Table 26)
BMI270 IMU	5 V	≈ 0.7 mA	Typical, full 6.4 kHz ODR (Bosch datasheet)
MG92B brake servo (each, 3×)	5 V	Not published	TowerPro does not publish a stall-current figure; measure during bring-up
MA600 encoder (each, 3×)	5 V	Not published	MPS datasheet does not state supply current in the accessible electrical table; measure during bring-up
moteus-n1 driver (each, 3×)	Pack	100 A peak / 9 A cont.	Datasheet (Table 26)
MN4006 motor (each, 3×, via driver)	Pack	16 A peak, thermally limited	Manufacturer bench data (Sec. 6.1)

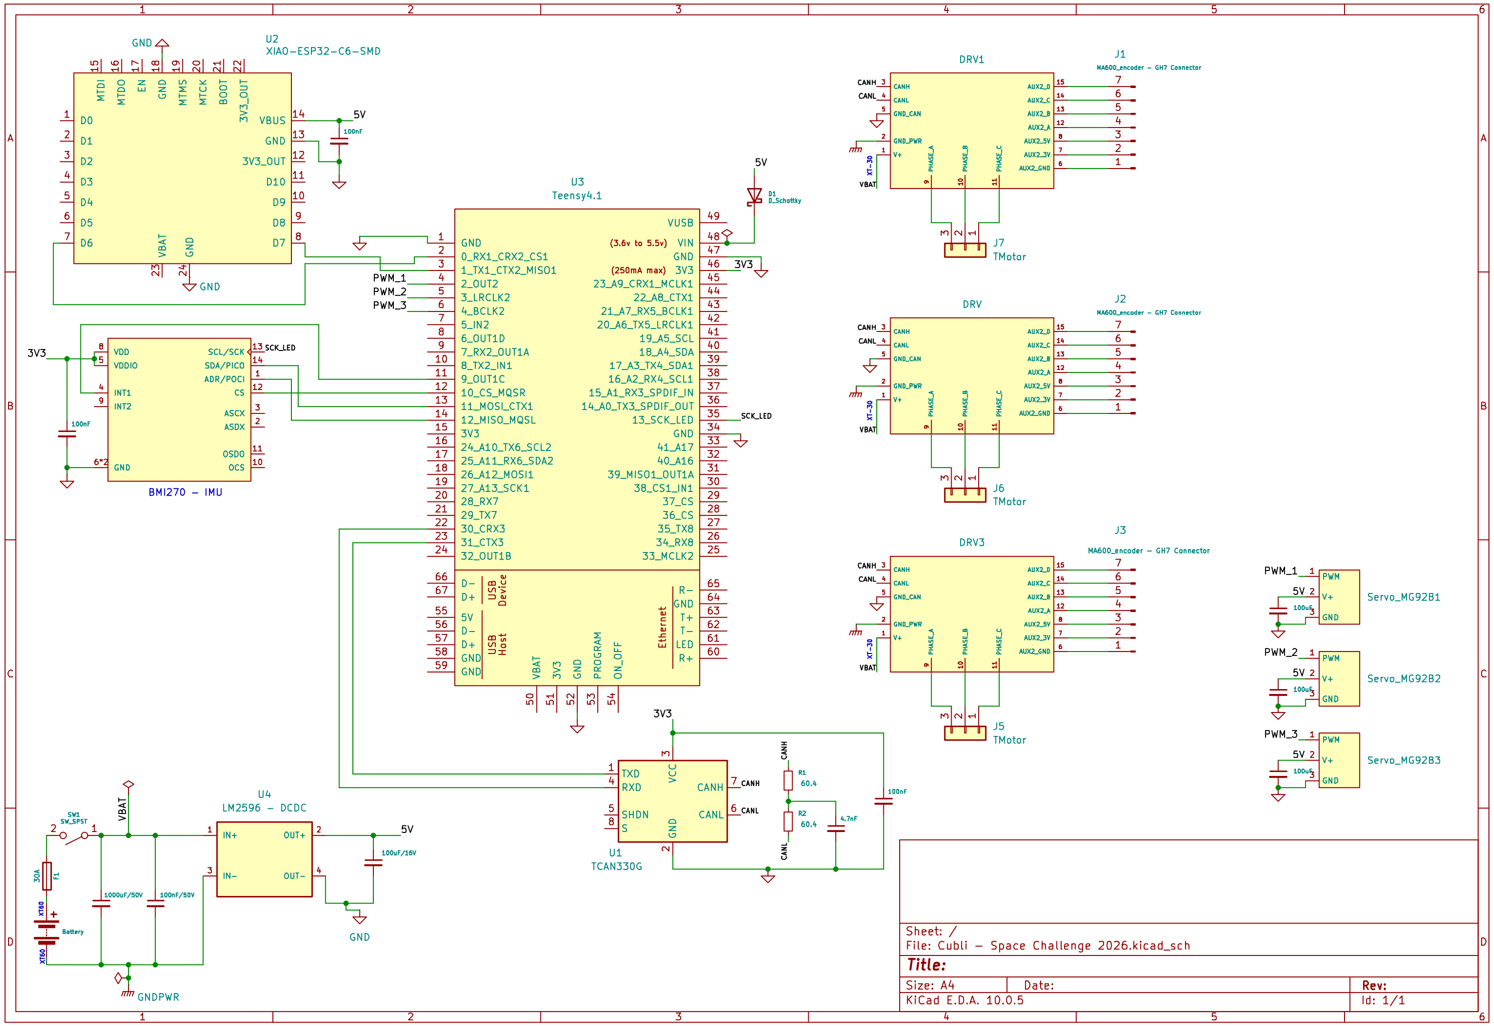

Figure 18: Full electrical schematic as captured, reproduced at page size. Covers the flight computer (Teensy 4.1) and its ESP32 telemetry module, the BMI270 IMU, the three moteus driver and motor connectors, the CAN-FD transceiver and termination, the LM2596 5 V converter from the pack, and the three brake servos.

C.3 Connector and Pin Assignments

Table 33 is the connector schedule. It exists so that a harness can be built and re-built from this document alone, and so that a mis-mate is a documented impossibility rather than a discovered one: no two connectors carrying different voltages share a housing type on this machine.

Table 33: Connector schedule. Keying and housing types are chosen so that two connectors carrying different voltages cannot be mated. Note: values have not been finalized at this preliminary stage.

Ref	From	To	Housing / keying	Signals
J1	Battery pack	Protection / E-stop	TBD	Pack +, pack −
J2	Protection	Distribution board	TBD	Motor bus
J3	Distribution board	Driver, axis 1	TBD	Motor bus
J4	Distribution board	Driver, axis 2	TBD	Motor bus
J5	Distribution board	Driver, axis 3	TBD	Motor bus
J6	Board	CAN chain, end A	TBD	CAN-H, CAN-L, GND
J7	CAN chain, end B	Termination network	TBD	CAN-H, CAN-L, GND
J8	Driver, axis i	Encoder, axis i	TBD	SPI, 3.3 V, GND
J9	Board	IMU	TBD	Serial bus, supply, GND
J10	Board	Telemetry module	TBD	UART, supply, GND
J11	Board	Brake servo, axis i	TBD	PWM, 5 V, GND

C.4 Harness Drawings and Wire Schedule

The harness is organised into seven runs: pack to protection, protection to board, board to the three driver drops, the three motor phase leads, the CAN chain segments, the three encoder SPI runs, and the three servo drives. Wire gauge on the motor bus (pack-to-protection, protection-to-board, driver drops and motor phase leads) is sized on spin-up current rather than the balancing current, since spin-up is the sustained high-current phase; the CAN and encoder SPI runs are sized on length limits rather than current (Table 17); the servo drives are sized on stall current. Gauge and run length are recorded once measured on the assembled harness.

C.5 Board Layout

The floorplan is constrained before it is aesthetic: the driver positions follow from the encoder cable limit, the bus route follows from the driver positions, and the board is placed around what remains (Section 6.5). Rail separation is maintained physically on the board, not only schematically, so that motor-current return does not share copper with the sensing ground.

C.6 Bring-Up and Isolation Checklists

These checks are performed on every re-assembly, not only on first build, since the encoder gaps and connector seating are exactly what mechanical work disturbs.

Before any power is applied.

1. Isolation: motor bus to chassis, logic rail to motor bus, each rail to ground.

2. Capacitor voltage ratings confirmed by position — no low-voltage part on the pack bus.
3. Polarity of every fabricated connector verified against Table 33.
4. Protection element in circuit: E-stop functional, fuse fitted and rated, anti-spark path intact.
5. Encoder air gaps re-measured after mechanical assembly.

First power-on, current-limited bench supply.

1. Current limit verified against a known load before connecting the assembly.
2. Logic rail voltage and ripple under load; converter thermal check.
3. All drivers enumerate on the bus; encoders read back with no error flags; IMU responds.
4. Bus integrity at the full data rate: sustained soak with all nodes addressed, error-frame count logged.
5. Open-loop wheel motion, wheels contained, one axis at a time.

First battery power-on. Performed only after the bench sequence passes in full, with the pack protection verified, cell monitoring active and the assembly in a fire-safe area. Inrush is measured against the bulk capacitance on connection.

C.7 As-Built Measurements

Table 34 records the measured electrical figures that replace the estimates used during design. These are the numbers reported back to the modelling and control work, and they are the reason the parameter identification of Section 3.3 runs on measured rather than assumed values.

Table 34: As-built electrical measurements. Each replaces an estimate used earlier in the design. Note: values have not been finalized at this preliminary stage.

Quantity	Measured	Replaces / feeds
Motor torque constant K_m	TBD	Table 7, datasheet value
Wheel inertia I_w (spin-down)	TBD	Table 7, CAD value
Viscous and Coulomb friction	TBD	C_b, C_w estimates
Command-to-encoder latency	TBD	Loop timing, Sec. 7.6.2
Achievable wheel speed on bus	TBD	$\omega_{\max}$, Sec. 4.1.1
Logic rail load, worst case	TBD	Table 32
MG92B servo stall current	TBD	Table 32, not published
MA600 encoder supply current	TBD	Table 32, not published
Runtime per charge under balance load	TBD	Demonstration planning

D Control Algorithm Detail